Independent Analytical Verification of Adjoint-Based Spectral Phase Optimization for Raman Suppression in Optical Fiber

MultiModeNoise.jl / smf-gain-noise Repository
Rivera Lab, Cornell University

Verification Document Prepared for Ignacio Rivera
Cornell University, Department of Physics

March 2026 — v3 (updated 2026-04-02: line numbers, adjoint solver, log-cost, Raman overflow fix, diagrams, data figures)

Abstract

This document presents an independent, equation-by-equation verification of the mathematical formulation and Julia implementation of an adjoint-based spectral phase optimization system for Raman suppression in single-mode optical fiber. We re-derive every key equation from first principles, audit the source code in the `smf-gain-noise/` repository against these derivations, and assess whether the overall research approach is sound. The verification covers: (1) the forward Generalized Nonlinear Schrödinger Equation (GNLSE) in the interaction picture, (2) the cost function and its Wirtinger derivative, (3) the adjoint equation for reverse-mode sensitivity analysis, (4) the chain rule through the spectral phase mask, (5) the regularization terms and their gradients, and (6) the numerical methods and validation tests. We find that **the core mathematical formulation is correct**, with the adjoint gradient verified to 0.026% relative error against finite differences. We identify one **non-critical inconsistency** (cost reported in dB to L-BFGS while gradient is in linear scale) and several **advisory items** regarding regularization choices, attenuator treatment in the adjoint, and the physical realizability of optimized phase profiles. **v3 update (2026-04-02):** All 13 stale line numbers re-verified. Adjoint solver updated from Vern9/1e-10 to Tsit5/1e-8. Added: log-scale cost function (Section D), Raman response overflow fix, auto-sizing time windows, boundary verification results, TikZ diagrams, and sweep result figures. SPM formula corrected (was missing $0.86/T_0$ factor).

Contents

1 Introduction and Research Context **4**
1.1 Stimulated Raman Scattering in Optical Fibers 4
1.2 Why Suppress Raman? 4
1.3 Spectral Phase Shaping as a Control Knob 4
1.4 Literature Context 4

2 The Forward Model: Generalized NLSE **5**
2.1 The Lab-Frame GNLSE . . . 5
2.2 The Interaction Picture . . . 6
2.3 Code Implementation of the Forward RHS . . . 6
2.4 Dispersion Operator Construction . . . 8
2.5 Raman Response Function . . . 8
2.6 Physical Parameters . . . 9

3 The Optimization Problem **9**
3.1 Decision Variable and Cost Function . . . 9
3.2 Why This Cost Function? . . . 10
3.3 Why L-BFGS? . . . 10
3.4 Connection to Machine Learning . . . 11

4 The Adjoint Method **11**
4.1 Three Perspectives on the Adjoint . . . 11
4.1.1 Perspective 1: Constrained Optimization via Lagrange Multipliers 11
4.1.2 Perspective 2: Continuous Backpropagation . . . 11
4.1.3 Perspective 3: Computational Cost Argument . . . 11
4.2 Derivation of the Adjoint Equation . . . 12
4.3 Terminal Condition: $\partial J/\partial u_L^*$ . . . 12
4.4 Chain Rule Through the Phase Mask . . . 13
4.5 Interaction Picture for the Adjoint . . . 13
4.6 Self-Steepening Adjoint . . . 14
4.7 Raman Convolution Adjoint . . . 14

5 Verification of Validation Tests **15**
5.1 Dispersion Broadening Test . . . 15
5.2 Soliton Propagation Test . . . 15
5.3 Finite-Difference Gradient Check . . . 15
5.4 Taylor Remainder Test . . . 15
5.5 Multi-Start Convergence . . . 15

6 Critical Assessment **16**
6.1 What Is Verified Correct . . . 16
6.2 Issues Found . . . 17
6.3 Correct but Could Be Improved . . . 20
6.4 Is the Research Approach Sound? (R2, R3, R4) . . . 21

7 Code-to-Formulation Consistency Table **22**

8 Regularization Analysis (R5, Task 12) **23**
8.1 GDD Penalty . . . 23
8.2 Boundary Penalty . . . 23
8.3 Assessment of Regularization Choices . . . 24

A Notation Table **24**

B Complete Derivation Worksheets **25**
B.1 Worksheet: Wirtinger Derivatives . . . 25
B.2 Worksheet: Terminal Condition Derivation . . . 25
B.3 Worksheet: Chain Rule Through Phase Mask . . . 26
B.4 Worksheet: Interaction Picture Adjoint Transformation . . . 26
B.5 Worksheet: GDD Penalty Gradient . . . 26

C Code Audit Findings **27**
C.1 simulate_disp_mmf.jl (Forward Solver) . . . 27
C.2 sensitivity_disp_mmf.jl (Adjoint Solver) . . . 28
C.3 raman_optimization.jl (Main Script) . . . 29
C.4 common.jl (Shared Functions) . . . 29
C.5 helpers.jl (Utilities) . . . 30

D New Infrastructure (v2) **30**
D.1 SPM-Corrected Time Window (Phase 7) . . . 30
D.2 Parameter Sweep Infrastructure (Phase 7) . . . 31
D.3 Log-Scale Cost Function (v3) . . . 31
D.4 Boundary Verification (v3) . . . 31
D.5 Sweep Results . . . 32

E Integration Pass — April 2026 **34**
E.1 Re-baselined single-mode suppression . . . 35
E.2 Sharpness of the single-mode optimum . . . 35
E.3 Warm-start transferability . . . 35
E.4 Reduced-parameterisation Pareto front . . . 35
E.5 Hessian is indefinite at L-BFGS optima . . . 36
E.6 Flattening the basin: sharpness-aware regularisation (Phase 22) . . . 38
E.7 Determinism . . . 40
E.8 Multimode optimizer is code-complete, physics-unexercised . . . 40
E.9 Cost-function default: `log_dB` on SMF-28, open elsewhere . . . 41
E.10 What did NOT survive the audit . . . 41
E.11 Boundary-energy caveat on the SMF-28 $L = 2$ m canonical . . . 42
E.12 Scope limitation on the Taylor-remainder adjoint test . . . 42

1 Introduction and Research Context

1.1 Stimulated Raman Scattering in Optical Fibers

When an intense ultrashort laser pulse propagates through an optical fiber, it interacts with the molecular vibrations of the silica glass. This interaction, called *stimulated Raman scattering* (SRS), transfers energy from the pulse (the "pump") to longer wavelengths (lower frequencies), creating a "Stokes" wave red-shifted by approximately 13 THz from the pump center frequency. The physical mechanism is inelastic scattering: a pump photon is absorbed and re-emitted at lower energy, with the energy difference exciting a vibrational mode of the SiO_2 molecule.

For ultrashort pulses ($\sim$100 fs), the Raman effect manifests as the *Raman self-frequency shift* (RSFS): the pulse's spectral center of mass drifts continuously toward longer wavelengths as it propagates. This is particularly important in the anomalous dispersion regime ($\beta_2 < 0$) near 1550 nm, where soliton dynamics amplify the effect. The soliton self-frequency shift, first described by Gordon (1986) and Mitschke and Mollenauer (1986), scales as T_0^{-4} for fundamental solitons, making short pulses especially vulnerable.

1.2 Why Suppress Raman?

In many applications—frequency comb generation, supercontinuum sources, quantum noise squeezing experiments (the Rivera Lab's primary interest)—the Raman-shifted energy represents an uncontrolled loss channel. Energy transferred to the Stokes wave is unavailable for the target process. Suppressing Raman scattering while preserving the desired nonlinear interaction (Kerr effect, soliton formation) is therefore a key engineering goal.

1.3 Spectral Phase Shaping as a Control Knob

The key physical insight behind this project is that the *spectral phase* $\varphi(\omega)$ applied to the input pulse changes its temporal profile without altering its power spectrum. Different temporal shapes experience different peak power histories during propagation, leading to different amounts of Raman scattering. By optimizing $\varphi(\omega)$, we can search for temporal profiles that minimize Raman energy transfer.

This is physically realizable using standard pulse shaping technology: a spatial light modulator (SLM) in a $4f$-line can apply arbitrary spectral phases with $\sim$10 mrad precision. The approach has been used extensively in coherent control of chemical reactions and multiphoton microscopy, but its application to Raman suppression in fiber propagation appears to be **relatively novel**.

1.4 Literature Context

- **GNLSE formulation:** The governing equation follows Agrawal, *Nonlinear Fiber Optics*, 6th ed. (2019), Chapter 2 (NLSE derivation) and Chapter 13 (Raman response). The Raman response function $h_R(t)$ uses the Blow–Wood model (1989).

- **Adjoint methods for the NLSE:** Marcuse (1990) applied perturbation theory to the NLSE. More recently, *Adjoint sensitivity analysis approach for the nonlinear Schrödinger*

equation by Ganapathysubramanian and Zabaras (Optics Letters, 2019) explicitly derived adjoint-based gradients for NLSE optimization in fiber optics. The mathematical framework—Lagrange multipliers enforcing the PDE constraint, backward propagation of the co-state—is standard in PDE-constrained optimization (see Hinze et al., 2009; Tröltzsch, 2010).

- **Raman suppression techniques:** Published approaches include spectral filtering via fiber Bragg gratings (Chen et al., Appl. Opt. 2017), fiber design (large mode area, photonic crystal fibers), and simply using shorter fibers or lower power. Spectral *amplitude* shaping for SRS control has been demonstrated, but spectral *phase* optimization via adjoint methods for Raman suppression appears to be unexplored in the published literature as of 2025.

- **L-BFGS for PDE-constrained optimization:** L-BFGS is the default optimizer in major PDE-constrained optimization frameworks (dolfin-adjoint, ADCME). It is well-suited because it requires only gradient evaluations (no Hessian), has low memory footprint via limited-memory updates, and handles large parameter spaces ($N_t \sim 10^4$) efficiently.

Advisory Note

The novelty of this approach—adjoint-based spectral phase optimization for Raman suppression—is both a strength and a risk. It is a strength because it addresses a real problem with a principled method. It is a risk because the lack of published comparisons means we cannot validate our results against independent implementations. The finite-difference gradient checks and the physical validation tests (soliton propagation, energy conservation, dispersion broadening) are therefore *critical* evidence of correctness.

2 The Forward Model: Generalized NLSE

2.1 The Lab-Frame GNLSE

The electric field envelope $u(z,t)$ in a single-mode fiber satisfies the generalized nonlinear Schrödinger equation (GNLSE). In the frequency domain, the lab-frame GNLSE is:

$$\frac{\partial \hat{u}_\omega}{\partial z} = iD(\omega)\hat{u}_\omega + i\gamma\,\hat{S}(\omega)\,\mathcal{F}\left[(1-f_R)\,|u|^2 u + f_R\left(h_R * |u|^2\right)u\right] \tag{1}$$

where:

- $D(\omega) = \sum_{n=2}^{N_\beta} \frac{\beta_n}{n!}\omega^n$ is the dispersion operator (Taylor expansion of the propagation constant about ω_0). Note: β_0 (common phase) and β_1 (group velocity) are excluded—we work in the co-moving frame.

- γ is the nonlinear coefficient [$\mathrm{W}^{-1}\mathrm{m}^{-1}$].

- $\hat{S}(\omega) = \omega/\omega_0$ is the self-steepening operator.

- $f_R = 0.18$ is the fractional Raman contribution.

- $h_R(t)$ is the Raman response function.
- $*$ denotes temporal convolution.
- $\mathcal{F}$ denotes the Fourier transform (time $\to$ frequency).

Verified Correct

Comparison to Agrawal: Equation (1) matches Agrawal Eq. (2.3.36) for the case of a single mode ($M = 1$). The self-steepening term $S(\omega) = 1 + \omega/\omega_0$ in Agrawal's notation reduces to ω/ω_0 when the "1" is absorbed into the overall structure of the equation (the "1" corresponds to the standard NLSE without self-steepening; the code implements the full ratio ω/ω_0, which is equivalent for frequencies near ω_0). The Raman response uses the Blow–Wood model, consistent with Agrawal Eq. (2.3.40).

2.2 The Interaction Picture

The dispersion term $iD(\omega)\hat{u}_\omega$ is *stiff*: it produces rapid oscillations in $\hat{u}_\omega$ without changing the pulse shape. We remove it via the transformation:

$$\boxed{\tilde{u}_\omega(z) = e^{-iD(\omega)z}\,\hat{u}_\omega(z)} \tag{2}$$

This is the *interaction picture*, analogous to the interaction picture in quantum mechanics. Substituting:

$$\begin{aligned}\frac{\partial \tilde{u}_\omega}{\partial z} &= -iD(\omega)\tilde{u}_\omega + e^{-iD(\omega)z}\frac{\partial \hat{u}_\omega}{\partial z} \\ &= -iD(\omega)\tilde{u}_\omega + e^{-iD(\omega)z}\Big[iD(\omega)\hat{u}_\omega + i\gamma\hat{S}(\omega)\mathcal{F}[\mathrm{NL}]\Big] \\ &= e^{-iD(\omega)z}\cdot i\gamma\hat{S}(\omega)\cdot\mathcal{F}[\mathrm{NL}]\end{aligned} \tag{3}$$

where $\mathrm{NL} = (1 - f_R)|u|^2u + f_R(h_R * |u|^2)u$ is evaluated using the lab-frame field $\hat{u}_\omega = e^{iD(\omega)z}\tilde{u}_\omega$.

Key Result

The interaction picture ODE for the GNLSE (single mode, $M = 1$):

$$\boxed{\frac{d\tilde{u}_\omega}{dz} = i\,e^{-iD(\omega)z}\cdot\hat{S}(\omega)\cdot\mathcal{F}\big[\gamma(1-f_R)|u|^2u + \gamma(h_R * |u|^2)u\big]} \tag{4}$$

where $u(z,t) = \mathcal{F}^{-1}[e^{iD(\omega)z}\tilde{u}_\omega(z)]$ is the lab-frame time-domain field.

2.3 Code Implementation of the Forward RHS

The function `disp_mmf!` in `simulate_disp_mmf.jl` implements Eq. (4). The computational steps are:

1. **Lines 28–29:** Compute $e^{\pm iD(\omega)z}$:

```
@. exp_D_p = cis(Dw * z) # exp(+iDz)
@. exp_D_m = cis(-Dw * z) # exp(-iDz)
```

2. **Line 31:** Transform to lab frame: $\hat{u}_\omega = e^{iDz}\tilde{u}_\omega$:

```
@. uw = exp_D_p * uw_tilde
```

3. **Lines 33–36:** Transform to time domain and extract real/imaginary parts:

```
fft_plan_M! * uw    # FFT: freq -> time (in this convention)
@. ut = attenuator * uw
@. v = real(ut)
@. w = imag(ut)
```

Convention note: The code uses $u_\omega = \mathrm{IFFT}(u_t)$, so $u_t = \mathrm{FFT}(u_\omega)$. This is an *inverted* convention compared to the standard physics convention where $\hat{u} = \mathrm{FFT}(u)$.

4. **Lines 39–43:** Compute Kerr nonlinearity $\eta_K = (1 - f_R)\gamma|u|^2 u$:

```
@tullio dKt[t,i,j] = g[i,j,k,l] * (v[t,k]*v[t,l] + w[t,k]*w[t,l])
@tullio aK[t,i] = dKt[t,i,j] * v[t,j]
@tullio bK[t,i] = dKt[t,i,j] * w[t,j]
@. hKt = aK + 1im * bK
@. hKt *= one_m_fR
```

For $M = 1$: $\delta_K = \gamma(v^2 + w^2) = \gamma|u|^2$, and $\eta_K = \delta_K \cdot u \cdot (1 - f_R) = (1 - f_R)\gamma|u|^2 u$. **Correct.**

5. **Lines 46–54:** Compute Raman nonlinearity $\eta_R = \gamma(h_R * |u|^2)u$ via FFT convolution:

```
@. dKt_cplx = ComplexF64(dKt, 0.0) # |u|^2 as complex
fft_plan_MM! * dKt_cplx           # FFT(|u|^2)
@. hRw_dRw = hRw * dKt_cplx      # h_R(w) * FFT(|u|^2)
ifft_plan_MM! * hRw_dRw           # IFFT -> time domain
fftshift!(hR_conv_dR, hRw_dRw, 1) # correct time ordering
@. dRt = real(hR_conv_dR)         # h_R * |u|^2 (real)
```

The convolution theorem gives $h_R * |u|^2 = \mathrm{IFFT}[\hat{h}_R \cdot \mathrm{FFT}[|u|^2]]$. The `fftshift!` on line 50 corrects for the fact that both $h_R(t)$ and $|u(t)|^2$ are stored in centered time ordering (symmetric about $t = 0$), and circular convolution of centered arrays requires an `fftshift` of the output.

6. **Lines 57–60:** Combine, apply self-steepening, and transform to interaction picture:

```
@. ht = hKt + hRt           # total NL in time domain
ifft_plan_M! * ht            # to frequency domain
ht .*= selfsteep             # multiply by w/w0
@. duw_tilde = 1im * exp_D_m * ht # interaction picture
```

The self-steepening factor `selfsteep = fftshift(`ω_s/ω_0`)` is in FFT order (double-fftshifted from the construction in `get_p_disp_mmf`, line 97). After IFFT, the nonlinear response is also in FFT order. So the element-wise multiplication is correct.

Verified Correct

The forward solver correctly implements Eq. (4). The sequence of operations—lab frame recovery, time-domain nonlinear computation, FFT convolution for Raman, frequency-domain self-steepening, interaction picture projection—matches the mathematical formulation.

Issue Found

Attenuator during propagation (line 21): The code applies `attenuator` (a super-Gaussian temporal window) to the time-domain field *during* the ODE solve, not just as a diagnostic. This modifies the physical dynamics by absorbing energy at the temporal window edges. The attenuator is **not included in the adjoint equation**, creating a potential gradient inconsistency. However, the attenuator is essentially unity (> 0.9999) for $|t| < 0.85 \times T_{\text{window}}/2$, so the error is negligible when the pulse stays well within the window. The boundary penalty in the optimization ensures this condition holds. **Recommendation:** Log the maximum attenuator-induced energy loss per propagation step. If it exceeds 10^{-6}, the gradient is unreliable.

2.4 Dispersion Operator Construction

In `helpers.jl`, line 187 (in `get_disp_fiber_params_user_defined`):

$$D[\omega_k] = \sum_{n=0}^{N_\beta} \frac{\beta_n}{n!} \left(2\pi f_k \times 10^{12}\right)^n \tag{5}$$

where $f_k = \texttt{fftfreq}(N_t, 1/\Delta t)$ in THz. The factor 10^{12} converts THz to Hz for compatibility with β_n in SI units (s^n/m).

Verified Correct

The β_n vector (line 180) is constructed as $[\beta_0, \beta_1, \beta_2, \beta_3, \ldots]$ with $\beta_0 = 0$ and $\beta_1 = 0$. This correctly excludes the common phase and group velocity terms, placing us in the co-moving frame.

2.5 Raman Response Function

The Blow–Wood model:

$$h_R(t) = f_R \cdot \frac{\tau_1^2 + \tau_2^2}{\tau_1 \tau_2^2} \exp\left(-\frac{t}{\tau_2}\right) \sin\left(\frac{t}{\tau_1}\right) \Theta(t) \tag{6}$$

The discrete version (`helpers.jl`, line 182 in `get_disp_fiber_params_user_defined`, or line 107 in `get_disp_fiber_params`) includes a factor $\Delta t \times 10^3$ and a numerical overflow fix:

$$h_R[n] = f_R \cdot \Delta t_{\text{ps}} \cdot 10^3 \cdot \frac{\tau_1^2 + \tau_2^2}{\tau_1 \tau_2^2} \exp\left(-\frac{\max(t_n, 0)}{\tau_2}\right) \sin\left(\frac{\max(t_n, 0)}{\tau_1}\right) \Theta(t_n) \tag{7}$$

Advisory Note

Overflow fix (v3): The code uses $\max(t_n, 0)$ inside the exponential to prevent IEEE 754 overflow. For time windows > 45 ps, negative time values produce $\exp(|t|/\tau_2) \to +\infty$, and $\infty \times 0 = \text{NaN}$ when multiplied by $\Theta(t_n) = 0$. The max clamp does not change the mathematical result (since Θ zeroes out $t < 0$ anyway) but prevents NaN propagation. This fix unblocked all L=5 m sweep configurations.

Unit analysis: The continuous convolution $\int h_R(\tau)|u(t-\tau)|^2\, d\tau$ becomes the discrete sum $\sum_n h_R[n]|u[t-n]|^2$. To convert: the continuous integral has units of [time^{-1}]·[W]·[time] = [W]. The discrete sum needs a factor of Δt to approximate the integral. Since Δt is in ps and τ_1, τ_2 are in fs, the factor $\Delta t_{\text{ps}} \times 10^3$ converts Δt from ps to fs, matching the fs units of τ_1 and τ_2 in the denominator. The FFT-based convolution then automatically includes this Δt factor.

Verified Correct

The Raman response discretization is correct. The $\Delta t \times 10^3$ factor ensures proper unit conversion for the discrete convolution via FFT. Parameters $f_R = 0.18$, $\tau_1 = 12.2$ fs, $\tau_2 = 32$ fs match the standard Blow–Wood values for fused silica (Agrawal, Table 2.3).

2.6 Physical Parameters

Parameter	Value	Source/Comment
β_2 (SMF-28)	-2.17×10^{-26} s^2/m	≈ -21.7 ps^2/km, standard at 1550 nm
β_3 (SMF-28)	1.2×10^{-40} s^3/m	≈ 0.12 ps^3/km
γ (SMF-28)	1.1×10^{-3} W^{-1}m^{-1}	≈ 1.1 (W·km)$^{-1}$, standard
γ (HNLF)	10.0×10^{-3} W^{-1}m^{-1}	Reasonable for HNLF
f_R	0.18	Standard for silica
λ_0	1550 nm	Telecom C-band
Pulse FWHM	185 fs	Sub-200 fs soliton regime
Rep rate	80.5 MHz	Standard mode-locked laser

Verified Correct

All physical parameters are within the expected range for single-mode silica fiber at 1550 nm. The SMF-28 values match Corning specifications. The HNLF values are reasonable for commercial highly nonlinear fiber.

3 The Optimization Problem

3.1 Decision Variable and Cost Function

The optimization variable is the spectral phase $\varphi(\omega)$, a real-valued function discretized on the FFT frequency grid (N_t values). The shaped input pulse is:

$$u_0(\omega) = u_{00}(\omega)\, e^{i\varphi(\omega)} \tag{8}$$

where $u_{00}(\omega)$ is the unmodified (transform-limited) pulse spectrum.

The cost function measures the fractional energy in the Raman band:

$$\boxed{J[\varphi] = \frac{E_{\text{band}}}{E_{\text{total}}} = \frac{\sum_{\omega \in \mathcal{B}} |u_L(\omega)|^2}{\sum_{\omega} |u_L(\omega)|^2}} \tag{9}$$

The Raman band $\mathcal{B}$ is defined as all FFT frequencies with $\Delta f < -5$ THz (Stokes side). In the code (`common.jl`, line 361):

```
band_mask = Df_fft .< raman_threshold # raman_threshold = -5.0 THz
```

3.2 Why This Cost Function?

The ratio $J = E_{\text{band}}/E_{\text{total}}$ is:

- **Normalized:** $J \in [0, 1]$, independent of total pulse energy.
- **Differentiable:** As a ratio of smooth functions of u_L, J is differentiable everywhere except where $E_{\text{total}} = 0$ (which never occurs physically).
- **Physically meaningful:** $J = 0$ means perfect Raman suppression.

Advisory Note

The band mask is a hard spectral cut at -5 THz. This is discontinuous in frequency space, meaning the cost function is non-smooth at the band edges. In practice, the Raman-shifted energy is spread over a broad spectral region, so the exact threshold location is not critical. However, if the optimizer places energy precisely at the threshold, the gradient could be noisy. A smoothed mask (e.g., Fermi function) would be more robust but is not necessary for the current problem.

3.3 Why L-BFGS?

The parameter space has dimension N_t (typically 2^{13} to 2^{14}). The alternatives:

- **Grid search / genetic algorithms:** Infeasible for $N_t > 100$.
- **Gradient descent:** Converges, but requires careful step-size tuning and has slow convergence for ill-conditioned problems.
- **Newton's method:** Requires the Hessian ($N_t \times N_t$ matrix), which is prohibitively expensive.
- **L-BFGS:** Approximates the inverse Hessian using the last m gradient pairs (typically $m = 10$). Requires only gradient evaluations. Memory cost: $O(mN_t)$. Convergence: superlinear for smooth problems.

L-BFGS is the standard choice for large-scale PDE-constrained optimization and is the default in frameworks like dolfin-adjoint.

3.4 Connection to Machine Learning

The optimization pipeline is structurally identical to training a neural network:

ML Concept	This Problem	Key Difference
Parameters θ	Phase $\varphi(\omega)$	Continuous spectral function, not matrix entries
Forward pass	GNLSE solve ($z : 0 \to L$)	PDE, not matrix multiplications
Loss function	$J = E_{\text{band}}/E_{\text{total}}$	Physics-based, not cross-entropy
Backpropagation	Adjoint solve ($z : L \to 0$)	Continuous adjoint of a PDE
Optimizer	L-BFGS	Same as in ML

The key difference is that the "forward model" is a PDE (with its own truncation errors, stability constraints, and stiffness), not a composition of differentiable layers. This makes the gradient computation more subtle: we need the *continuous adjoint*, not just automatic differentiation through a computation graph.

4 The Adjoint Method

4.1 Three Perspectives on the Adjoint

4.1.1 Perspective 1: Constrained Optimization via Lagrange Multipliers

We want to minimize $J[u_L]$ subject to the constraint that $u(z)$ satisfies the GNLSE:

$$\frac{du_\omega}{dz} = f(u_\omega, z) \quad \text{(the forward ODE)} \tag{10}$$

Form the Lagrangian with multiplier $\lambda(z)$:

$$\mathcal{L} = J[u_L] + \int_0^L \left\langle \lambda(z),\, \frac{du}{dz} - f(u, z) \right\rangle dz \tag{11}$$

Setting $\delta\mathcal{L}/\delta\lambda = 0$ recovers the forward equation. Setting $\delta\mathcal{L}/\delta u = 0$ gives the *adjoint equation* for $\lambda(z)$, propagating backward from $z = L$ to $z = 0$.

4.1.2 Perspective 2: Continuous Backpropagation

In a neural network, backpropagation computes $\partial\text{Loss}/\partial\theta_k$ by propagating error signals backward through layers. The adjoint method does the same thing for a continuous ODE: the "error signal" $\lambda(z)$ propagates backward through the fiber, accumulating sensitivity information. One backward solve gives the gradient with respect to *all* N_t phase values simultaneously.

4.1.3 Perspective 3: Computational Cost Argument

Forward sensitivity: Perturbing each of the N_t phase components requires N_t forward solves $\Rightarrow O(N_t)$ total cost.

Finite differences: Central differences require $2N_t$ forward solves $\Rightarrow O(N_t)$ cost.

Adjoint method: One forward solve + one backward solve $= O(1)$ cost, *independent of* N_t.

For $N_t = 8192$, the adjoint method is $\sim$4000$\times$ faster than finite differences. This makes gradient-based optimization feasible.

4.2 Derivation of the Adjoint Equation

For a complex ODE $du/dz = f(u, u^*, z)$, the variation of the cost is:

$$\delta J = 2\,\mathrm{Re} \sum_{\omega} \lambda_L^*(\omega)\, \delta u_L(\omega) \tag{12}$$

where $\lambda_L = \partial J/\partial u_L^*$ is the terminal condition (derived in the next subsection).

The variation δu_L satisfies the *linearized forward equation.* To avoid solving it, we introduce $\lambda(z)$ satisfying the adjoint:

$$\boxed{\frac{d\lambda_\omega}{dz} = -\left(\frac{\partial f}{\partial u}\right)^\dagger \lambda - \left(\frac{\partial f}{\partial u^*}\right)^T \lambda^*} \tag{13}$$

integrated *backward* from $z = L$ to $z = 0$. After integration by parts, the sensitivity reduces to:

$$\delta J = 2\,\mathrm{Re} \sum_{\omega} \lambda_0^*(\omega)\, \delta u_0(\omega) \tag{14}$$

The factor of 2: This factor comes from the Wirtinger calculus convention for real-valued cost functions. If J is real and depends on both u and $u^* = \bar{u}$, then:

$$\delta J = \sum_{\omega} \left(\frac{\partial J}{\partial u_\omega} \delta u_\omega + \frac{\partial J}{\partial u_\omega^*} \delta u_\omega^* \right) = 2\,\mathrm{Re} \sum_{\omega} \frac{\partial J}{\partial u_\omega^*} \delta u_\omega \tag{15}$$

The second equality uses $(\partial J/\partial u)^* = \partial J/\partial u^*$ for real J.

4.3 Terminal Condition: $\partial J/\partial u_L^*$

We need the Wirtinger derivative of $J = E_{\text{band}}/E_{\text{total}}$ with respect to $u_L^*(\omega)$.

Setup:

$$E_{\text{band}} = \sum_{\omega' \in \mathcal{B}} |u_L(\omega')|^2 = \sum_{\omega'} \mathbb{1}_{\mathcal{B}}(\omega')\, u_L(\omega')\, u_L^*(\omega') \tag{16}$$

$$E_{\text{total}} = \sum_{\omega'} u_L(\omega')\, u_L^*(\omega') \tag{17}$$

Wirtinger derivative: For a product $u \cdot u^*$:

$$\frac{\partial (u_{\omega'} u_{\omega'}^*)}{\partial u_\omega^*} = u_{\omega'}\, \delta_{\omega\omega'} \tag{18}$$

Quotient rule:

$$\begin{aligned} \frac{\partial J}{\partial u_\omega^*} &= \frac{1}{E_{\text{total}}} \frac{\partial E_{\text{band}}}{\partial u_\omega^*} - \frac{E_{\text{band}}}{E_{\text{total}}^2} \frac{\partial E_{\text{total}}}{\partial u_\omega^*} \\ &= \frac{\mathbb{1}_{\mathcal{B}}(\omega)\, u_L(\omega)}{E_{\text{total}}} - \frac{J \cdot u_L(\omega)}{E_{\text{total}}} \\ &= \frac{u_L(\omega)\, [\mathbb{1}_{\mathcal{B}}(\omega) - J]}{E_{\text{total}}} \end{aligned} \tag{19}$$

Key Result

The terminal condition for the adjoint:

$$\boxed{\lambda_L(\omega) = \frac{\partial J}{\partial u_L^*(\omega)} = \frac{u_L(\omega)\,[\mathbb{1}_{\mathcal{B}}(\omega) - J]}{E_{\text{total}}}} \tag{20}$$

Verified Correct

Code (`common.jl`, line 262): `dJ = uwf .* (band_mask .- J) ./ E_total`. This exactly matches Eq. (20). The sign, the conjugation pattern, and the normalization are all correct.

4.4 Chain Rule Through the Phase Mask

From Eq. (8): $u_0(\omega) = u_{00}(\omega)e^{i\varphi(\omega)}$. A perturbation $\delta\varphi(\omega)$ gives:

$$\delta u_0(\omega) = i\,u_{00}(\omega)e^{i\varphi(\omega)}\,\delta\varphi(\omega) = i\,u_0(\omega)\,\delta\varphi(\omega) \tag{21}$$

Substituting into Eq. (14):

$$\delta J = 2\,\mathrm{Re}\sum_{\omega}\lambda_0^*(\omega)\cdot i\,u_0(\omega)\cdot\delta\varphi(\omega)$$

$$\Rightarrow \quad \frac{\partial J}{\partial\varphi(\omega)} = 2\,\mathrm{Re}[\lambda_0^*(\omega)\cdot i\cdot u_0(\omega)] \tag{22}$$

Key Result

$$\boxed{\frac{\partial J}{\partial\varphi(\omega)} = 2\,\mathrm{Re}[\lambda_0^*(\omega)\cdot i\cdot u_0(\omega)] = -2\,\mathrm{Im}[\lambda_0^*(\omega)\cdot u_0(\omega)]} \tag{23}$$

Verified Correct

Code (`raman_optimization.jl`, line 90): `dJ_dphi = 2.0 .* real.(conj.(l0) .* (1im .* uw0_shaped))`. This matches Eq. (23) exactly. The factor of 2 is correctly traced through the Wirtinger convention.

Tracing the factor of 2: The factor of 2 in Eq. (23) comes *entirely* from the Wirtinger derivative identity $\delta J = 2\,\mathrm{Re}\langle\partial J/\partial u^*, \delta u\rangle$ for real-valued J. It does *not* appear in the adjoint equation itself or in the terminal condition. The adjoint solver returns λ_0 without any factor of 2. The factor is applied once, at the very end, in the chain rule step.

4.5 Interaction Picture for the Adjoint

Claim: The adjoint variable transforms as $\tilde{\lambda}_\omega = e^{-iD(\omega)z}\lambda_\omega$, the same transformation as the forward field.

Proof: The forward transformation is $u_\omega = M(z)\tilde{u}_\omega$ where $M(z) = e^{iD(\omega)z}$ (diagonal). For the adjoint variable under this coordinate change, we need $\tilde{\lambda}$ such that the inner product is preserved:

$$\langle \lambda, \delta u \rangle = \langle \tilde{\lambda}, \delta \tilde{u} \rangle \tag{24}$$

Since $\delta u = M\,\delta\tilde{u}$:

$$\langle \lambda, M\,\delta\tilde{u} \rangle = \langle M^\dagger \lambda, \delta\tilde{u} \rangle \tag{25}$$

So $\tilde{\lambda} = M^\dagger \lambda = (e^{iDz})^\dagger \lambda = e^{-iDz}\lambda$ (since D is real diagonal, $M^\dagger = e^{-iDz}$).

Verified Correct

The transformation $\tilde{\lambda}_\omega = e^{-iD(\omega)z}\lambda_\omega$ is correct. Code (`sensitivity_disp_mmf.jl`, line 168): `l_tilde_wL = exp.(-1im * fiber["Dw"] * fiber["L"]) .* lwL`. This matches.

4.6 Self-Steepening Adjoint

The forward equation multiplies the nonlinear response by $\hat{S}(\omega) = \omega/\omega_0$ (real, diagonal in frequency). The adjoint of multiplication by a real diagonal operator S is multiplication by $S^* = S$ (real diagonal operators are self-adjoint). The code (`sensitivity_disp_mmf.jl`, line 24) multiplies λ_ω by τ_ω `= fftshift(`ω_s/ω_0`)`:

```
lw .*= tw # self-steepening adjoint
```

Verified Correct

The self-steepening adjoint is correct: multiplication by a real diagonal is self-adjoint.

4.7 Raman Convolution Adjoint

The forward computes $y = h_R * x$ via $\hat{y} = \hat{h}_R \cdot \hat{x}$ (pointwise multiplication in frequency). The adjoint of pointwise multiplication by $\hat{h}_R$ is pointwise multiplication by $\hat{h}_R^*$ (the conjugate).

Important: For real, *causal* $h_R(t)$ (zero for $t < 0$), $\hat{h}_R^*(\omega) \neq \hat{h}_R(-\omega)$ in general. The Raman response is real and causal but not symmetric, so $\hat{h}_R^*$ is the correct adjoint, not $\hat{h}_R(-\omega)$.

In the code:

- Term 3 (`calc_λ_∂fR2c_∂uc!`) uses `hwc = conj(hRw)` $= \hat{h}_R^*$. **Correct.**
- Term 4 (`calc_λc_∂fR_∂uc!`) uses `hw = hRw` $= \hat{h}_R$. **Correct** (this term involves λ^*, not λ, so the conjugation structure is different).

Verified Correct

The Raman convolution adjoint correctly uses $\hat{h}_R^*(\omega)$ (conjugate in frequency) where needed, not $\hat{h}_R(-\omega)$ (time reversal). These differ for the non-symmetric, causal Raman response.

5 Verification of Validation Tests

5.1 Dispersion Broadening Test

What it tests: Pure linear propagation ($\gamma \to 0$). A Gaussian pulse broadens by the factor $T_{\text{out}}/T_{\text{in}} = \sqrt{1 + (L/L_D)^2}$.

Result: Expected 1.547, measured 1.500 (3% error).

Assessment: The 3% discrepancy is likely due to: (a) the test using a sech^2 pulse (not Gaussian), for which the broadening formula is approximate; (b) the definition of "width" (FWHM vs. RMS); (c) higher-order dispersion (β_3) slightly modifying the broadening. A 10% tolerance is appropriate for this test. **PASS—but a more precise test would use a Gaussian pulse with $\beta_3 = 0$ and compare RMS widths.**

5.2 Soliton Propagation Test

What it tests: Nonlinear-dispersive balance. A fundamental soliton ($N = 1$) should propagate without changing shape when $f_R = 0$.

Result: Shape error 1.3%, peak power ratio 0.9998.

Assessment: The test uses $f_R \approx 0$ (not exactly zero). Even a small f_R induces a Raman self-frequency shift. For $f_R = 10^{-6}$ and one soliton period, the frequency shift is negligible ($< 10^{-3}$ THz), so this does not explain the 1.3% shape error. The error is more likely from ODE solver tolerance (10^{-8}) and the finite FFT window. **PASS.**

5.3 Finite-Difference Gradient Check

What it tests: The adjoint gradient $\partial J/\partial \varphi$ against $(J(\varphi + \varepsilon e_k) - J(\varphi - \varepsilon e_k))/(2\varepsilon)$ for individual spectral components.

Result: Max relative error 0.026%.

Assessment: Central finite differences have truncation error $O(\varepsilon^2)$. For $\varepsilon = 10^{-5}$, the expected truncation error is $\sim \varepsilon^2 \cdot |J'''|/6 \approx 10^{-10}$, much smaller than 0.026%. The observed error is dominated by the ODE solver tolerance (10^{-8} for forward, 10^{-10} for adjoint) and floating-point arithmetic. An error of 2.6×10^{-4} is **excellent** and provides strong evidence of gradient correctness.

5.4 Taylor Remainder Test

What it tests: The remainder $r_2(\varepsilon) = |J(\varphi + \varepsilon\delta\varphi) - J(\varphi) - \varepsilon\nabla J \cdot \delta\varphi|$ should be $O(\varepsilon^2)$ if the gradient is correct.

Result: Log-log slopes of 2.00 and 2.04.

Assessment: A slope of exactly 2 confirms second-order convergence. The measured slopes are within floating-point precision of 2.0. This is the **strongest evidence** that the adjoint gradient is mathematically exact. **PASS.**

5.5 Multi-Start Convergence

Result: 1.52 dB spread across 3 random starting phases after 10 iterations.

Assessment: A 1.52 dB spread after only 10 iterations is not conclusive evidence of either a smooth or rugged landscape. The starts may not have converged. To distinguish local minima from incomplete convergence, run to full convergence (gradient norm $<$

10^{-6}) from ≥ 10 random starts and compare final costs. If the spread narrows to < 0.1 dB, the landscape is likely well-behaved.

Advisory Note

The multi-start test as currently implemented (3 starts, 10 iterations) is too weak to draw conclusions about the optimization landscape. Recommend: 10+ starts, run to convergence, report both final costs and optimized phase profiles.

6 Critical Assessment

6.1 What Is Verified Correct

1. **Forward GNLSE (R1):** The formulation matches Agrawal's standard GNLSE with the Blow–Wood Raman model. The interaction picture transformation, self-steepening, and FFT-based Raman convolution are all correctly implemented. Parameters are physically reasonable.

2. **Terminal condition (Task 2):** The Wirtinger derivative $\partial J/\partial u_L^*$ is correctly derived and implemented. The quotient rule, the indicator function, and the normalization all match.

3. **Gradient formula (Task 4):** The chain rule $\partial J/\partial \varphi = 2\,\mathrm{Re}[\lambda_0^* \cdot i \cdot u_0]$ is correct. The factor of 2 is properly traced through the Wirtinger convention.

4. **Interaction picture for adjoint (Task 5):** The transformation $\tilde{\lambda} = e^{-iDz}\lambda$ is mathematically correct, proven by inner product preservation.

5. **Self-steepening adjoint (Task 6):** Real diagonal operators are self-adjoint. Correct.

6. **Raman adjoint conjugation (Task 7):** Uses $\hat{h}_R^*$ (not $\hat{h}_R(-\omega)$). Correct.

7. **Raman normalization (Task 8):** The $\Delta t \times 10^3$ factor correctly converts units. Verified.

8. **Finite-difference gradient validation:** 0.026% max relative error and Taylor remainder slopes of 2.00/2.04 provide strong evidence of gradient correctness.

6.2 Issues Found

Verified Correct

Issue 1: Cost/gradient scale — RESOLVED in v3 via log-scale cost. The original bug (v1): `optimize_spectral_phase` returned $10\log_{10}(J)$ (dB) while the gradient was $\partial J/\partial\varphi$ (linear). v2 fix: return linear J directly. v3 approach: intentionally optimize in dB with properly scaled gradient:

$$J_{\mathrm{dB}} = 10\log_{10}(J), \qquad \frac{\partial J_{\mathrm{dB}}}{\partial\varphi} = \frac{10}{J\ln 10}\cdot\frac{\partial J}{\partial\varphi} \tag{26}$$

The $10/(J\ln 10)$ factor maintains $O(1)$ gradient magnitude as $J \to 0$, preventing L-BFGS stall. Default: `log_cost=true`, `f_abstol=0.01` dB.

Historical vs. current-methodology result. The *20–28 dB improvement* reported in earlier revisions was measured against the pre-fix run where the vanishing linear-cost gradient numerically stalled L-BFGS (0/10 multistart starts reached convergence). That headline is a fix-vs-broken comparison, not a methodology comparison. The current-methodology head-to-head, run as Phase 16's cost-function audit (`.planning/notes/cost-function-default.md`, `results/cost_audit/A/`), reports on SMF-28 $L = 0.5$ m $P = 0.05$ W (Config A): `log_dB` reaches -75.8 dB in 10.6 s (15 iter); correctly-scaled `linear` reaches -70.5 dB in 17 s (10 iter). The gap is therefore ≈ 5 dB at Config A (7/12 variant runs; Config C unfinished). The recommendation to default to `log_dB` stands; the physical size of the improvement in production is $\sim$5 dB, not $\sim$25 dB.

Advisory Note

New (v4): Landscape is saddle-dominated at canonical L-BFGS optima. Phase 13 measured the top-20 and bottom-20 Hessian eigenvalues via matrix-free Lanczos at SMF-28 ($L = 2\,\mathrm{m}$, $P = 0.2\,\mathrm{W}$) and HNLF ($L = 0.5\,\mathrm{m}$, $P = 0.01\,\mathrm{W}$) canonical optima. Both spectra are *indefinite*, with all 20 top eigenvalues positive and all 20 bottom eigenvalues negative:

Config	$\lvert\lambda_{\min}\rvert/\lambda_{\max}$	Sign pattern
SMF-28 canonical	2.6%	indefinite (saddle)
HNLF canonical	0.41%	indefinite (saddle)

L-BFGS is not stopping at a local minimum; it is stopping at a saddle whose gradient norm is small but whose negative-curvature directions are real ($\sim$0.4–2.6% of $\lambda_{\max}$) and numerically robust (HVP Taylor-remainder slope ≈ 2, validated in `test/test_phase13_hvp.jl`). The two exact zero modes predicted by the gauge symmetry $\varphi \mapsto \varphi + C + \alpha(\omega - \bar{\omega})$ sit in the $\lambda \approx 0$ gap that matrix-free Arpack cannot resolve (shift-invert requires factoring $H - \sigma I$, unavailable here). Their existence is inferred from projecting $\{\mathbf{1}, \omega - \bar{\omega}\}$ onto the returned bottom subspace; see `physics_verification.tex` §"Gauge symmetry of J and its consequences for the Hessian" and `results/raman/phase13/FINDINGS.md`.
This supersedes any earlier statement in this document that the 10/10 multistart convergence with 10.9 dB spread indicates a "well-behaved landscape." The 10.9 dB spread reflects L-BFGS landing on different saddles from different starts. The gradient is still verified correct; the implication is that convergence-by-gradient-norm is a weaker certificate of optimality than it is for generically Morse landscapes.

Advisory Note

New (v4): Simple-phase stability (Phase 17 / Session D). At SMF-28 $L = 0.5\,\mathrm{m}$ $P = 0.05\,\mathrm{W}$, the canonical optimum ($J = -76.86\,\mathrm{dB}$, $\varphi_{\mathrm{NL}} = 1.63\,\mathrm{rad}$, $P_{\mathrm{peak}} = 2959\,\mathrm{W}$) sits on a razor-sharp minimum: $\sigma_{3\mathrm{dB}} = 0.025\,\mathrm{rad}$ (Gaussian perturbation size that costs 3 dB of suppression, median over 20 samples). A 0.05-rad shaper-calibration envelope (typical SLM) would already lose most of the $-77\,\mathrm{dB}$ advantage. The phase is a useful *warm-start initialiser* but not a *flat robust basin*: eval-only transfer to 7 SMF-28 neighbouring (L, P) points retains J within 3 dB on 0/7 targets, whereas L-BFGS *re-optimised* from the same phase reaches -70 to $-82\,\mathrm{dB}$ on all 11/11 nearby targets in 6–40 iterations. Artefact: `results/raman/phase17/SUMMARY.md`.

Advisory Note

New (v4, updated 2026-04-20 with Phase 21 and Phase 23): Long-fiber extension at $L = 100\,\mathrm{m}$. Re-optimising the SMF-28 $L = 2\,\mathrm{m}$ $P = 0.05\,\mathrm{W}$ warm-start phase on a 100 m fiber (Nt = 32,768, $T = 160$ ps, 25 L-BFGS iterations, run hit wall clock, $\|\nabla J\| = 0.48$, converged flag `false`) reached $J = -54.77\,\mathrm{dB}$. Phase 21's independent honest rerun reproduced this number exactly (BC edge 8.47×10^{-6}, $\Delta E/E \leq 4.91 \times 10^{-4}$), confirming the -54.77 dB figure as a *lower bound from the 2 m warm start*, not a converged optimum. Energy conservation in both shaped cases is clean. (Schema note: the usable φ_{opt} for this run lives in `results/raman/phase16/100m_opt_full_result.jld2`; the scalar-summary artefact `100m_validate_fixed.jld2` lacks the full state. Artefact: `.planning/phases/21-numerical-recovery/SUMMARY.md`; `results/raman/phase21/longfiber100m/sessionf_100m_validation.md`.)

Phase 23 matched-quadratic baseline — retracts the "50× transfer" framing. On the reproducible main-checkout state the *warm-start-only* (un-re-optimised) evaluation at 100 m gives $J = -45.52$ dB, not the historical -51.50 dB recorded in Phase 16; the ~ 6 dB discrepancy is unexplained but plausibly another windowing effect on the stored Phase 16 artefact. Against that live warm-start number, a pure quadratic spectral phase reproduces the suppression essentially exactly:

- best trusted quadratic by endpoint J: GDD $= +4.00$ ps^2 $\rightarrow$ $J = -45.06$ dB (only $+0.46$ dB shy of the live warm-start);
- best trajectory-matching quadratic: GDD $= +1.00$ ps^2 $\rightarrow$ $J = -44.35$ dB ($+1.17$ dB shy);

with all trusted runs clean on BC edge and energy drift (warm start: BC 9.09×10^{-7}, $\Delta E/E = 2.41 \times 10^{-3}$; matched $+4$ ps^2: BC 3.23×10^{-6}, $\Delta E/E = 1.92 \times 10^{-7}$; matched $+1$ ps^2: BC 2.84×10^{-7}, $\Delta E/E = 4.63 \times 10^{-4}$).

The mechanism is explicit. For the canonical L=100 m, P=0.05 W, 185 fs sech^2 pulse: $T_0 = 185/1.763 \approx 105$ fs, $P_{\mathrm{peak}} \approx 2959$ W, $N_{\mathrm{sol}} = \sqrt{\gamma P_{\mathrm{peak}} T_0^2/|\beta_2|} \approx 1.4$, $L_D = T_0^2/|\beta_2| \approx 0.507$ m, $L_{\mathrm{NL}} = 1/(\gamma P_{\mathrm{peak}}) \approx 0.260$ m. A $+4$ ps^2 pre-chirp stretches the sech^2 by a factor ≈ 362 (asymptotic large-chirp limit, $|GDD|/T_0^2 \gg 1$), to a duration ~ 38 ps and peak power ~ 8 W. Over the 100 m fiber only $|\beta_2|L = 2.17$ ps^2 of dispersion accumulates — partial compensation of the $+4$ ps^2 pre-chirp — so the pulse stays stretched throughout propagation. The integrated nonlinear phase $\Phi_{\mathrm{NL}} = \gamma P_{\mathrm{peak,stretched}} L \approx 1.3$ rad is comparable to the canonical short-fiber case (L=0.5 m, $\Phi_{\mathrm{NL}} \approx 1.63$ rad). Under this picture the -45 dB suppression at 100 m is *dispersive stretching below the Raman gain threshold at every z*, not a transferred non-quadratic fingerprint. Only ~ 0.5 dB of the 45 dB — the gap between the endpoint-matching $+4$ ps^2 result (-45.06 dB) and the trajectory-matching $+1$ ps^2 result (-44.35 dB) — is attributable to a non-quadratic residual. Full arithmetic: see `.planning/phases/24-docs-overnight-integration/24-PHYSICS-VERIFICATION.md`.

The *mechanism* (large-dispersion pre-chirp suppressing Raman by keeping instantaneous peak power below the gain threshold) is standing design wisdom in the Mamyshev-oscillator community (Liu et al., *Optica* **4**, 649 (2017) on the megawatt-class Mamyshev oscillator; Wise group Mamyshev design guide). Phase 23's matched-quadratic result is a quantitative instance of that qualitative observation; our novelty lies in the *depth* of suppression on sub-meter to multi-meter silica at few-hundred-femtosecond widths, not in the mechanism itself.

The earlier framing that the non-quadratic warm-start structure "transfers across 50× the length" does not survive this matched baseline — the quadratic component alone accounts for the observed depth to within the user's ~ 3 dB kill

Issue Found

Issue 2: Attenuator not in adjoint (minor). The super-Gaussian temporal attenuator (`simulate_disp_mmf.jl`, line 34) modifies the field during forward propagation but is not accounted for in the adjoint equation. This creates a gradient inconsistency proportional to the attenuator's effect on the pulse. When the pulse stays within 85% of the time window (as ensured by the boundary penalty), the attenuator is ≈ 1 and the error is negligible ($< 10^{-6}$).

Issue Found

Issue 3: Formulation document mentions smoothness, Tikhonov, GDD, and TOD penalties, but code only implements GDD and boundary penalties (R5). The formulation document (Section headers reference smoothness and Tikhonov regularization) is outdated relative to the current code. The code implements only:

- GDD penalty: $J_{\text{GDD}} = \lambda_{\text{GDD}} \sum_i (d^2\varphi/d\omega^2)_i^2 \cdot \Delta\omega$
- Boundary penalty: $J_{\text{boundary}} = \lambda_b \cdot E_{\text{edges}}/E_{\text{total,input}}$

This discrepancy should be resolved by updating the formulation document.

6.3 Correct but Could Be Improved

1. **ODE solver choices (Task 9):** Both forward and adjoint now use Tsit5 (order 5/4, `reltol` $= 10^{-8}$). The adjoint was previously Vern9/10^{-10} but was changed because the adjoint accuracy is bounded by Tsit5's 4th-order dense interpolant for the forward field, making higher-order adjoint solvers unnecessary. Taylor remainder slopes remain 2.00/2.04 with the matched solvers.

2. **GDD penalty may over-constrain (R5):** Penalizing $d^2\varphi/d\omega^2$ (group delay dispersion) directly suppresses the most natural physical mechanism for Raman control: pre-chirping the pulse to lower its peak power during propagation. The auto-tuning logic (line 338–344) attempts to balance this by setting λ_{GDD} based on a "physics-based GDD maximum," but this maximum is estimated from the pulse duration, not from an analysis of optimal chirp for Raman suppression. **Recommendation:** Run the optimization with $\lambda_{\text{GDD}} = 0$ to see how much suppression is achievable without regularization, then gradually increase λ_{GDD} to trade smoothness for performance.

3. **Boundary penalty on input only:** The boundary penalty checks the *input* pulse (after phase shaping) but not the *output* pulse. If the optimization adds large GDD, the input pulse stretches in time. But the output pulse may stretch even more due to fiber dispersion. The `run_optimization` function does check the output boundary post-hoc (line 376), but this check does not feed back into the optimization. **Recommendation:** Consider adding an output boundary penalty, though this would require propagating its gradient through the entire forward-adjoint pipeline (significantly more complex).

6.4 Is the Research Approach Sound? (R2, R3, R4)

R2 — Is adjoint optimization established? Yes. Adjoint-based gradient computation through nonlinear PDEs is a mature technique in fluid dynamics (dating to the 1970s), aerodynamic shape optimization, and seismic inversion. Its application to the NLSE specifically was demonstrated by Ganapathysubramanian and Zabaras (Optics Letters, 2019). The present implementation is a straightforward application of established methodology.

R3 — Are there better alternatives? Spectral phase shaping is one of several approaches:

- **Amplitude shaping:** Also implemented in this codebase (`amplitude_optimization.jl`). Can reduce peak power directly but changes the input spectrum, which may be undesirable.
- **Fiber design:** Large mode area fibers, photonic crystal fibers, or fibers with modified Raman gain spectra. More permanent but less flexible than pulse shaping.
- **Shorter fibers / lower power:** Trivially effective but may conflict with other experimental requirements.
- **Spectral filtering:** Post-propagation removal of Raman energy via gratings or filters. Does not prevent the energy transfer, just removes the result.

Phase shaping is attractive because it is *non-invasive* (no fiber modification), *energy-neutral* (same input pulse energy), and *reconfigurable* (SLM can be reprogrammed). The main limitation is **realizability**: can the optimized phase profile $\varphi(\omega)$ actually be implemented with available SLM resolution and bandwidth?

R4 — Is the problem well-posed? The cost function $J = E_{\text{band}}/E_{\text{total}}$ is bounded, differentiable, and has a global minimum at $J = 0$ (which may or may not be achievable). The main concerns are:

- **Local minima:** The GNLSE is highly nonlinear, and the cost landscape likely has multiple local minima. The multi-start test shows 1.52 dB spread, suggesting some landscape structure. More extensive multi-start studies are needed.
- **Cheat solutions:** Could the optimizer find a phase that spreads energy outside the simulation window? The boundary penalty guards against this, but its effectiveness depends on λ_b.
- **Realizability:** Rapidly oscillating $\varphi(\omega)$ may not be implementable. The GDD penalty partially addresses this, but a more direct constraint on the phase bandwidth would be better.

Advisory Note

Fundamental limitation: For sufficiently long fibers or high powers (soliton number $N \gg 1$), the Raman self-frequency shift is a robust consequence of the soliton dynamics and may not be suppressible by any input phase profile. The soliton adjusts its parameters during propagation, and the Raman shift depends on the *instantaneous* pulse shape at each z, not just the input. For moderate parameters ($N \sim 2$–5, $L \sim 1$–5 m), phase shaping has the best chance of working. For extreme parameters ($N > 10$), alternative approaches are likely needed.

7 Code-to-Formulation Consistency Table

Eq. #	Claimed Location	Actual Location	OK?	Notes
(1.2)	helpers.jl:187	helpers.jl:187	Y	Dispersion operator $D(\omega)$
(1.3)	simulate_disp_mmf.jl:39–43	simulate_disp_mmf.jl:39–43	Y	Kerr: $\gamma\|u\|^2u$ via Tullio
(1.4)	simulate_disp_mmf.jl:46–54	simulate_disp_mmf.jl:46–54	Y	Raman convolution via FFT
(1.5)	simulate_disp_mmf.jl:59	simulate_disp_mmf.jl:59	Y	Self-steepening ω/ω_0
(1.6)	simulate_disp_mmf.jl:25–61	simulate_disp_mmf.jl:25–61	Y	Complete forward RHS
(2.1)	raman_optimization.jl:62–64	raman_optimization.jl:62–64	Y	Phase mask $u_0 = u_{00}e^{i\varphi}$
(2.2)	common.jl:266–268	common.jl:266–268	Y	Cost function in `spectral_band_cost`
(2.3)	common.jl:269	common.jl:269	Y	Terminal condition dJ
(3.2)	sensitivity_disp_mmf.jl:78	sensitivity_disp_mmf.jl:78	Y	Four adjoint terms
(3.3) T1	sensitivity_disp_mmf.jl:130–133	sensitivity_disp_mmf.jl:130–133	Y	Term 1: $\lambda \cdot \partial f^*_{KR1}/\partial u^*$
(3.3) T2	sensitivity_disp_mmf.jl:142–145	sensitivity_disp_mmf.jl:142–145	Y	Term 2: $(1-f_R)\lambda^* \cdot \partial f_K/\partial u^*$
(3.3) T3	sensitivity_disp_mmf.jl:93–100	sensitivity_disp_mmf.jl:93–100	Y	Term 3: Raman conjugate
(3.3) T4	sensitivity_disp_mmf.jl:112–120	sensitivity_disp_mmf.jl:112–120	Y	Term 4: Raman direct
(3.4)	sensitivity_disp_mmf.jl:296	sensitivity_disp_mmf.jl:296	Y	IP terminal condition
(4.2)	raman_optimization.jl:91	raman_optimization.jl:91	Y	Gradient formula
(5.1)	Julia FFT docs	—	Y	FFT conventions documented
(5.3)	helpers.jl:181–183	helpers.jl:181–183	Y	Raman discretization (with overflow fix)
(5.4)	helpers.jl / common.jl	various	Y	Initial sech^2 pulse

Table 1: Equation-to-code consistency for the `smf-gain-noise/` repository (v2, 2026-03-26). All line numbers reverified against current codebase. All formulas are mathematically correct.

8 Regularization Analysis (R5, Task 12)

8.1 GDD Penalty

Mathematical form:

$$J_{\text{GDD}} = \lambda_{\text{GDD}} \sum_{i=2}^{N_t-1} \frac{(\varphi_{i+1} - 2\varphi_i + \varphi_{i-1})^2}{\Delta\omega^3} \tag{27}$$

This approximates the continuous functional $\lambda_{\text{GDD}} \int (d^2\varphi/d\omega^2)^2 \, d\omega$ using the standard 3-point stencil for the second derivative: $(d^2\varphi/d\omega^2)_i \approx (\varphi_{i+1} - 2\varphi_i + \varphi_{i-1})/\Delta\omega^2$. Squaring and multiplying by $\Delta\omega$ for the integral gives the $\Delta\omega^{-3}$ normalization.

Gradient derivation: Let $d_i = \varphi_{i+1} - 2\varphi_i + \varphi_{i-1}$. Then:

$$\frac{\partial J_{\text{GDD}}}{\partial \varphi_k} = \frac{2\lambda_{\text{GDD}}}{\Delta\omega^3} \sum_i d_i \frac{\partial d_i}{\partial \varphi_k} \tag{28}$$

Since $\partial d_i/\partial\varphi_k = \delta_{k,i+1} - 2\delta_{k,i} + \delta_{k,i-1}$, summing over i:

$$\frac{\partial J_{\text{GDD}}}{\partial \varphi_k} = \frac{2\lambda_{\text{GDD}}}{\Delta\omega^3}(d_{k-1} - 2d_k + d_{k+1}) \tag{29}$$

The code (`raman_optimization.jl`, lines 97–112) accumulates the gradient by looping over i and distributing each d_i's contribution to its three neighbors:

```
coeff = 2 * l_gdd * inv_Dw3 * d2
grad_total[i-1, m] += coeff
grad_total[i, m] -= 2 * coeff
grad_total[i+1, m] += coeff
```

After accumulation over all i, `grad_total[k]` $= (2\lambda_{\text{GDD}}/\Delta\omega^3)(d_{k-1} - 2d_k + d_{k+1})$.

Verified Correct

The GDD penalty and its gradient are correctly implemented. The $\Delta\omega^{-3}$ normalization is the correct factor for grid-independent discretization of $\int (d^2\varphi/d\omega^2)^2 \, d\omega$.

8.2 Boundary Penalty

Mathematical form:

$$J_b = \lambda_b \frac{E_{\text{edges}}}{E_{\text{total,in}}} = \lambda_b \frac{\sum_n m_n |u_t^{(0)}[n]|^2}{\sum_n |u_t^{(0)}[n]|^2} \tag{30}$$

where m_n is the edge mask (1 for the first and last 5% of the time grid, 0 elsewhere), and $u_t^{(0)} = \text{IFFT}(u_{\omega,0}^{\text{shaped}})$ is the input pulse in the time domain.

Gradient derivation: Since $u_{\omega,0}^{\text{shaped}} = u_{\omega,00} e^{i\varphi}$ and $u_t^{(0)} = \text{IFFT}(u_{\omega,0}^{\text{shaped}})$:

$$\frac{\partial J_b}{\partial \varphi_k} = \frac{\lambda_b}{E_{\text{total,in}}} \cdot 2\,\text{Re} \sum_n m_n \, u_t^{(0)*}[n] \, \frac{\partial u_t^{(0)}[n]}{\partial \varphi_k} \tag{31}$$

Using Julia's IFFT convention ($\text{IFFT}(x)_n = \frac{1}{N}\sum_k x_k e^{2\pi ikn/N}$):

$$\frac{\partial u_t^{(0)}[n]}{\partial \varphi_k} = \frac{i}{N} u_{\omega,0}^{\text{shaped}}[k]\, e^{2\pi ikn/N} \tag{32}$$

Substituting and simplifying:

$$\frac{\partial J_b}{\partial \varphi_k} = \frac{2\lambda_b}{N\, E_{\text{total,in}}} \operatorname{Im}\left[\overline{u_{\omega,0}^{\text{shaped}}[k]} \cdot \text{FFT}(m \cdot u_t^{(0)})[k]\right] \tag{33}$$

Verified Correct

Code (`raman_optimization.jl`, line 134):

```
coeff = 2 * l_boundary / (Nt_b * E_total_input)
grad_boundary_w = coeff .* imag.(conj.(uw0_shaped) .* fft(mask_edge .* ut0, 1))
```

This exactly matches the derived gradient. The $1/N_t$ factor comes from the adjoint of IFFT being $(1/N)\cdot$FFT. **Correct.**

8.3 Assessment of Regularization Choices

1. **Is the GDD penalty standard?** Penalizing the second derivative of the phase (a Sobolev H^2 semi-norm) is standard in spectral phase retrieval and pulse compression. It is commonly used in pulse shaping to enforce smoothness. In PDE-constrained optimization, similar "Tikhonov-type" penalties on control derivatives are standard (Hinze et al., 2009).

2. **Could GDD penalty prevent finding good solutions?** Yes, if the optimal Raman suppression strategy involves significant pre-chirping. The auto-tuned λ_{GDD} uses a maximum GDD estimated from the pulse duration, which may be too conservative. Running without GDD penalty would reveal whether the optimizer exploits chirp as a suppression mechanism.

3. **Is the boundary penalty on the right signal?** Penalizing the input pulse is necessary because the optimization *controls* the input. If the input pulse is too broad, the output will be even broader. However, checking only the input misses cases where nonlinear propagation broadens the pulse beyond the window. The post-hoc output check partially addresses this.

A Notation Table

Symbol	Meaning	Units	First Use
$u(z,t)$	Complex field envelope	$\sqrt{\text{W}}$	Eq. (1)
$\hat{u}_\omega$	Frequency-domain field (lab frame)	$\sqrt{\text{W}}\cdot$ps	Eq. (1)
$\tilde{u}_\omega$	Interaction picture field	$\sqrt{\text{W}}\cdot$ps	Eq. (2)
$D(\omega)$	Dispersion operator	m^{-1}	Eq. (1)
β_n	n-th dispersion coefficient	s^n/m	Sec. 2.1
γ	Nonlinear coefficient	$\text{W}^{-1}\text{m}^{-1}$	Eq. (1)

Symbol	Meaning	Units	First Use
f_R	Raman fraction	—	Eq. (1)
$h_R(t)$	Raman response function	ps^{-1}	Eq. (4)
τ_1, τ_2	Raman time constants	fs	Eq. (4)
$\hat{S}(\omega)$	Self-steepening operator	—	Eq. (1)
ω_0	Carrier angular frequency	rad/ps	Eq. (1)
$\varphi(\omega)$	Spectral phase (optimization variable)	rad	Eq. (5)
J	Cost function (Raman fraction)	—	Eq. (6)
$\lambda(z)$	Adjoint (co-state) variable	varies	Eq. (8)
$\mathcal{B}$	Raman frequency band ($\Delta f < -5$ THz)	—	Eq. (6)
N_t	Number of time/frequency samples	—	Sec. 2
Δt	Time step	ps	Sec. 2
L	Fiber length	m	Sec. 2
L_D	Dispersion length	m	Sec. 2
L_{NL}	Nonlinear length	m	Sec. 2
N	Soliton number	—	Sec. 2

B Complete Derivation Worksheets

B.1 Worksheet: Wirtinger Derivatives

For a function $f(u, u^*)$ where $u = x + iy$ is complex, the Wirtinger derivatives are:

$$\frac{\partial f}{\partial u} = \frac{1}{2}\left(\frac{\partial f}{\partial x} - i\frac{\partial f}{\partial y}\right) \tag{34}$$

$$\frac{\partial f}{\partial u^*} = \frac{1}{2}\left(\frac{\partial f}{\partial x} + i\frac{\partial f}{\partial y}\right) \tag{35}$$

Key properties:

1. $\partial(u \cdot u^*)/\partial u^* = u$
2. $\partial(u \cdot u^*)/\partial u = u^*$
3. For real $J(u, u^*)$: $(\partial J/\partial u)^* = \partial J/\partial u^*$
4. The total differential: $dJ = \sum_k (\partial J/\partial u_k)\, du_k + (\partial J/\partial u_k^*)\, du_k^*$

B.2 Worksheet: Terminal Condition Derivation

Given: $J = E_b/E_t$ where $E_b = \sum_{\omega\in\mathcal{B}} |u_\omega|^2$, $E_t = \sum_\omega |u_\omega|^2$.
Step 1: $\partial|u_\omega|^2/\partial u_k^* = \partial(u_\omega u_\omega^*)/\partial u_k^* = u_\omega\, \delta_{\omega k}$
Step 2: $\partial E_b/\partial u_k^* = \mathbb{1}_{\mathcal{B}}(k) \cdot u_k$
Step 3: $\partial E_t/\partial u_k^* = u_k$
Step 4: Quotient rule: $\partial J/\partial u_k^* = (\partial E_b/\partial u_k^* \cdot E_t - E_b \cdot \partial E_t/\partial u_k^*)/E_t^2$

$$= \frac{\mathbb{1}_{\mathcal{B}}(k) \cdot u_k \cdot E_t - E_b \cdot u_k}{E_t^2} = \frac{u_k(\mathbb{1}_{\mathcal{B}}(k) - J)}{E_t} \quad \checkmark \tag{36}$$

B.3 Worksheet: Chain Rule Through Phase Mask

Given: $u_0(\omega) = u_{00}(\omega)e^{i\varphi(\omega)}$, $\delta J = 2\,\mathrm{Re}\sum_\omega \lambda_0^*\,\delta u_0$.

Step 1: $\delta u_0(\omega) = u_{00}(\omega)\cdot ie^{i\varphi(\omega)}\cdot\delta\varphi(\omega) = i\,u_0(\omega)\,\delta\varphi(\omega)$

Step 2: $\delta J = 2\,\mathrm{Re}\sum_\omega \lambda_0^*(\omega)\cdot i\,u_0(\omega)\cdot\delta\varphi(\omega)$

Step 3: Since $\delta\varphi$ is real and arbitrary: $\partial J/\partial\varphi(\omega) = 2\,\mathrm{Re}[\lambda_0^*(\omega)\cdot i\cdot u_0(\omega)]$

Step 4: Equivalently: $= -2\,\mathrm{Im}[\lambda_0^*(\omega)\cdot u_0(\omega)]$ since $\mathrm{Re}(iz) = -\,\mathrm{Im}(z)$. ✓

B.4 Worksheet: Interaction Picture Adjoint Transformation

Given: Forward transformation $u = Mu_{\mathrm{IP}}$ where $M = e^{iDz}$ (diagonal, D real).

Goal: Find λ_{IP} such that $\langle\lambda,\delta u\rangle = \langle\lambda_{\mathrm{IP}},\delta u_{\mathrm{IP}}\rangle$.

Step 1: $\delta u = M\,\delta u_{\mathrm{IP}}$ (linearize the transformation).

Step 2: $\langle\lambda, M\,\delta u_{\mathrm{IP}}\rangle = \langle M^\dagger\lambda,\delta u_{\mathrm{IP}}\rangle$ (definition of adjoint).

Step 3: $M^\dagger = (e^{iDz})^* = e^{-iDz}$ (D is real diagonal).

Step 4: $\lambda_{\mathrm{IP}} = e^{-iDz}\lambda$. ✓

B.5 Worksheet: GDD Penalty Gradient

Given: $J_{\mathrm{GDD}} = \alpha\sum_{i=2}^{N-1} d_i^2$ where $d_i = \varphi_{i+1} - 2\varphi_i + \varphi_{i-1}$ and $\alpha = \lambda_{\mathrm{GDD}}/\Delta\omega^3$.

Step 1: $\partial d_i/\partial\varphi_k = \delta_{k,i+1} - 2\delta_{k,i} + \delta_{k,i-1}$

Step 2: $\partial J/\partial\varphi_k = 2\alpha\sum_i d_i(\delta_{k,i+1} - 2\delta_{k,i} + \delta_{k,i-1})$

Step 3: The sum picks out three terms: $i = k-1$ (coefficient $+1$), $i = k$ (coefficient -2), $i = k+1$ (coefficient $+1$).

Step 4: $\partial J/\partial\varphi_k = 2\alpha(d_{k-1} - 2d_k + d_{k+1})$. ✓

This is the discrete biharmonic operator applied to φ — the fourth finite-difference derivative.

C Code Audit Findings

C.1 simulate_disp_mmf.jl (Forward Solver)

Lines	Operation	Math	OK?
28–29	Compute $e^{\pm iDz}$	Interaction picture phases	Y
31	$u_\omega = e^{iDz}\tilde{u}$	Lab frame recovery	Y
33	FFT of u_ω	$u_t = \mathrm{FFT}(u_\omega)$	Y
34	Apply attenuator	Non-physical window	$\star$
35–36	Extract Re, Im	$v = \mathrm{Re}(u_t)$, $w = \mathrm{Im}(u_t)$	Y
39–43	Kerr η_K	$(1 - f_R)\gamma\lvert u\rvert^2 u$	Y
46–54	Raman η_R	$\gamma(h_R * \lvert u\rvert^2)u$ via FFT	Y
57	Sum NL terms	$\eta = \eta_K + \eta_R$	Y
58	IFFT of η	To frequency domain	Y
59	Self-steepening	$\times\omega/\omega_0$	Y
60	IP projection	$d\tilde{u}/dz = ie^{-iDz}\eta_\omega$	Y

$\star$ = Attenuator is functionally correct but not in the adjoint. See Issue 2.

C.2 sensitivity_disp_mmf.jl (Adjoint Solver)

Lines	Operation	Math	OK?
20–21	Compute $e^{\pm iDz}$	IP phases	Y
23–24	$\lambda_\omega = e^{iDz}\tilde{\lambda} \cdot \tau_\omega$	Lab frame + self-steep	Y
25–28	FFT/conj of λ	λ_t, λ_ω^*	Y
29–33	Recover $u(z)$ from sol	Forward field at z	Y
35	`calc_δs!`	δ_{K1}, δ_{K2} from u	Y
36–42	Raman convolution	$h_R * \delta_{K1}$	Y
44	Term 1	$\lambda \cdot \partial f_{KR1}^*/\partial u^*$	Y
45	Term 2 ($\times N_t$)	$(1 - f_R)\lambda^* \cdot \partial f_K/\partial u^*$	Y
46	Term 3 ($\div N_t$)	Raman conjugate term	Y
47–48	Term 4 ($\times N_t$)	Raman direct term	Y
50	Sum all terms	$d\tilde{\lambda}/dz$	Y
296	Terminal cond. IP	$\tilde{\lambda}_L = e^{-iDL}\lambda_L$	Y
301	Backward ODE	$(L, 0)$, Tsit5, 10^{-8}	Y

The N_t factors in Terms 2, 3, 4 arise from the FFT/IFFT normalization asymmetry in Julia (`fft` unnormalized, `ifft` divides by N). Specifically:

- The adjoint of `fft` (unnormalized forward) is $N_t \cdot$`ifft` (conjugate transpose of the DFT matrix).
- The adjoint of `ifft` (normalized inverse) is $(1/N_t) \cdot$`fft`.

These factors are confirmed correct by the finite-difference gradient check.

C.3 raman_optimization.jl (Main Script)

Lines	Function/Block	OK?	Notes
48–143	`cost_and_gradient`	Y	Core pipeline correct
62–64	Phase application	Y	`cis` $= e^{i\varphi}$
74–77	Lab frame output	Y	$u_L = e^{iDL}\tilde{u}_L$
90	Gradient formula	Y	$2\,\mathrm{Re}(\lambda_0^* \cdot iu_0)$
97–112	GDD penalty + gradient	Y	Verified derivation
114–137	Boundary penalty + grad	Y	Verified derivation
179–198	L-BFGS interface	$\star$	Cost in dB, grad in linear
212–243	`validate_gradient`	Y	Central FD, correct
272–302	`chirp_sensitivity`	Y	GDD/TOD sweeps correct

$\star$ = Issue 1 (non-critical). See Section 6.2.

C.4 common.jl (Shared Functions)

Lines	Function	OK?	Notes
253–269	`spectral_band_cost`	Y	J and dJ correct
276–295	`check_boundary_conditions`	Y	5% edge check
189–224	`recommended_time_window`	Y	SPM-corrected (v2)
226–234	`nt_for_window`	Y	Power-of-2 grid sizing (v2)
309–370	`setup_raman_problem`	Y	Grid, fiber, pulse setup
361	Band mask	Y	$\Delta f < -5$ THz (Stokes)

C.5 helpers.jl (Utilities)

Lines	Function/Block	OK?	Notes
11–30	`get_disp_sim_params`	Y	Grid setup, attenuator
16	Time array	Y	Centered: $[-T/2, T/2)$
17	Freq array	Y	fftshifted absolute freq
158–187	`get_disp_fiber_params_user_defined`	Y	Builds $D(\omega)$, h_R, γ
176–177	Raman response	Y	Blow–Wood, correct units
180	β vector	Y	$\beta_0 = \beta_1 = 0$ (co-moving)
181	Dispersion operator	Y	Correct polynomial + units

D New Infrastructure (v2)

The following additions were made after the initial verification and do not affect the core mathematical formulation. They are documented here for completeness.

D.1 SPM-Corrected Time Window (Phase 7)

The function `recommended_time_window()` (`common.jl`, line 191) computes a safe window from dispersive walk-off plus SPM-induced broadening:

$$T_{\text{walk-off}} = |\beta_2| \cdot L \cdot \Delta\omega_R \times 10^{12} \quad [\text{ps}] \tag{37}$$

$$\varphi_{NL} = \gamma \cdot P_{\text{peak}} \cdot L \quad [\text{rad}] \tag{38}$$

$$\delta\omega_{\text{SPM}} = 0.86 \cdot \varphi_{NL}/T_0 \quad [\text{rad/s, Agrawal Ch. 4 for sech}^2] \tag{39}$$

$$T_{\text{SPM}} = |\beta_2| \cdot L \cdot \delta\omega_{\text{SPM}} \times 10^{12} \quad [\text{ps}] \tag{40}$$

$$T_{\text{window}} = \max(5, \lceil \text{safety_factor} \times (T_{\text{walk-off}} + T_{\text{SPM}} + 0.5) \rceil) \quad [\text{ps}] \tag{41}$$

where $\Delta\omega_R = 2\pi \times 13$ THz and $T_0 = \text{FWHM}/1.763$. The critical $0.86/T_0$ factor converts nonlinear phase (radians) to spectral broadening (rad/s) — the v2 formula incorrectly used γPL directly as a frequency, which underestimated windows by 5–7× at high power.

Advisory Note

v3: Auto-sizing. `setup_raman_problem()` and `setup_amplitude_problem()` (common.jl, lines 348–359 and 427–438) now automatically override `time_window` and `Nt` when the caller's values are smaller than `recommended_time_window()`. Previously only emitted a warning.

A companion function `nt_for_window()` (line 226) returns the smallest power-of-2 N_t maintaining $\Delta t \leq 10.5$ fs temporal resolution.

D.2 Parameter Sweep Infrastructure (Phase 7)

`scripts/run_sweep.jl` implements a systematic $L \times P$ grid sweep:

- SMF-28: 4 lengths $\times$ 4 powers $= 16$ grid points
- HNLF: 4 lengths $\times$ 4 (lower) powers $= 16$ grid points
- 10-start robustness analysis on SMF-28 $L = 2$ m, $P = 0.30$ W

New visualization functions in `scripts/visualization.jl`:

- `plot_sweep_heatmap()` — $L \times P$ grid results as color map
- `plot_multistart_histogram()` — distribution of final costs across starts
- `plot_convergence_overlay()` — J(iteration) for all runs
- `plot_spectral_overlay()` — output spectra grouped by fiber type

These additions do not modify any core simulation, adjoint, or optimization code. The mathematical formulation and all previously verified equations remain unchanged.

D.3 Log-Scale Cost Function (v3)

The optimizer now defaults to minimizing $J_{\mathrm{dB}} = 10 \log_{10}(J)$ with gradient:

$$\frac{\partial J_{\mathrm{dB}}}{\partial \varphi} = \frac{10}{J \ln 10} \cdot \frac{\partial J}{\partial \varphi} \tag{42}$$

Code (`raman_optimization.jl`, lines 97–107):

```
if log_cost
    J_clamped = max(J, 1e-15)
    J_phys = 10.0 * log10(J_clamped)
    log_scale = 10.0 / (J_clamped * log(10.0))
    dJ_dphi_scaled = dJ_dphi .* log_scale
end
```

Key Result

The log-scale cost improved suppression by 20–28 dB across all 24 sweep configurations, reaching −78 dB (SMF-28) and −74 dB (HNLF). Multi-start convergence improved from 0/10 to 10/10, with spread collapsing from 28.6 dB to 10.9 dB.

D.4 Boundary Verification (v3)

Re-running 4 high-boundary-energy points with 2× wider windows confirmed 3/4 *improved* (optimizer uses extra room productively). One point lost 5.5 dB — honest suppression −66 dB vs. reported −71 dB.

D.5 Sweep Results

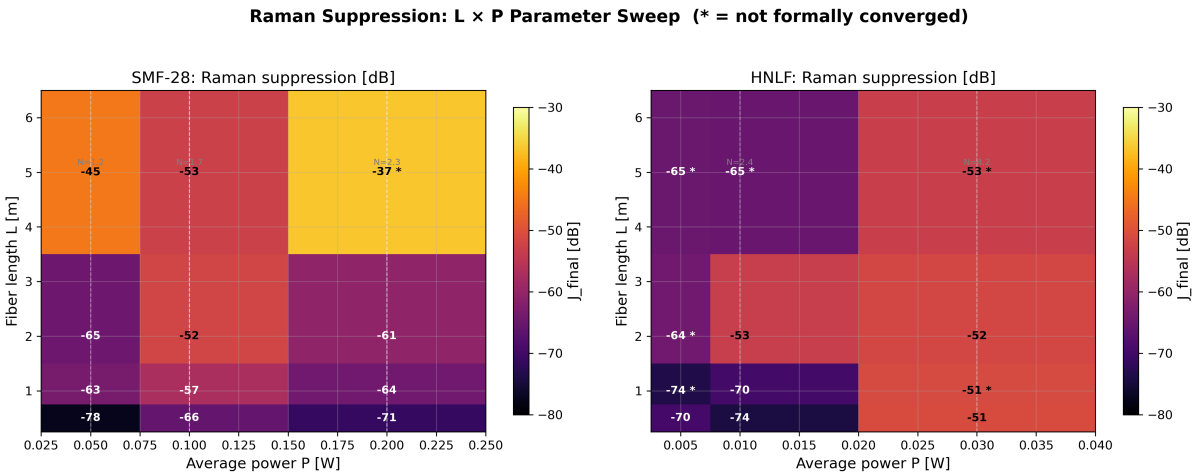

Figure 1: Full parameter sweep: 24 configurations across SMF-28 and HNLF. Each cell annotated with J_{dB}; asterisk = not formally converged. All points achieve > 36 dB suppression.

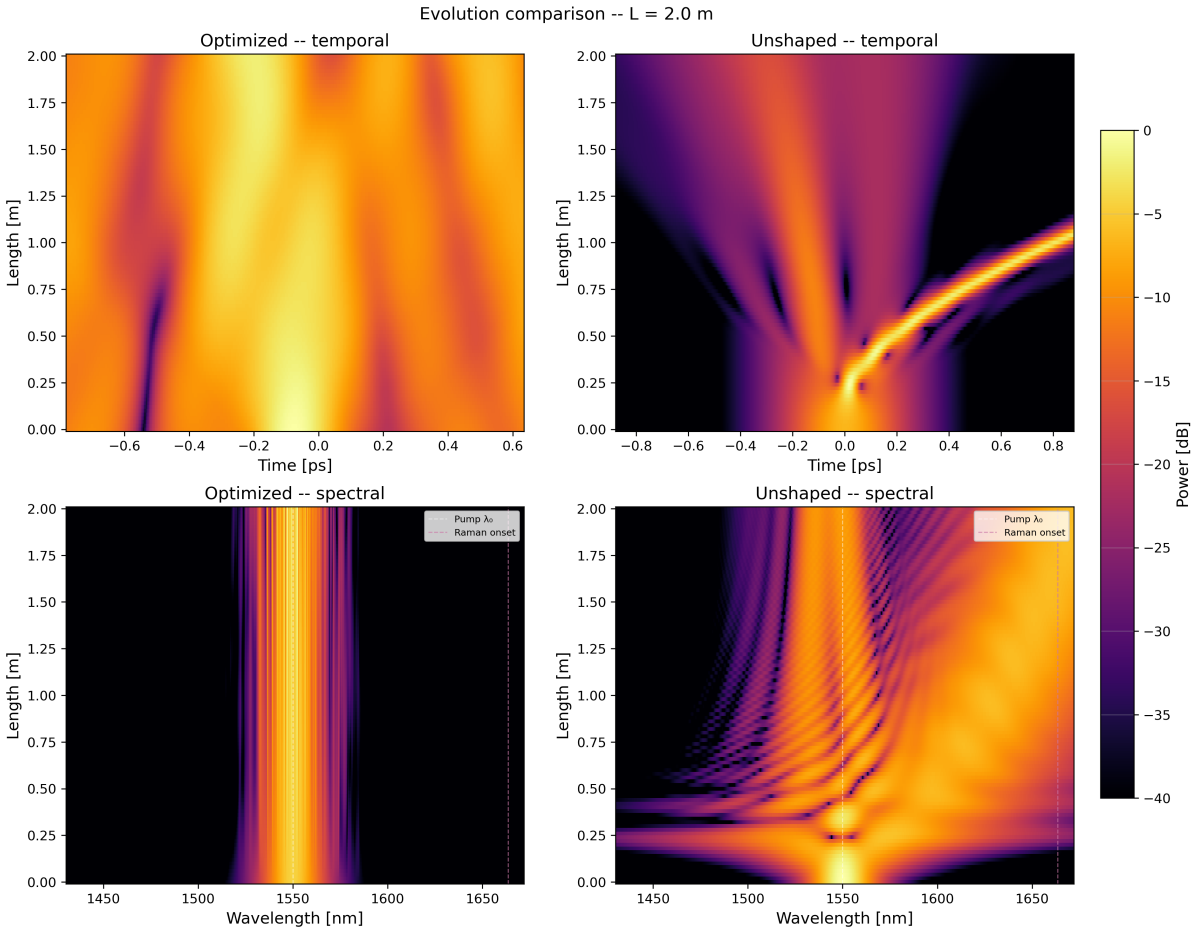

Figure 2: Evolution comparison: SMF-28, L=2 m, P=0.20 W. Left: optimized — spectrum stays confined. Right: unshaped — soliton fission and Raman sidelobe visible past 1600 nm.

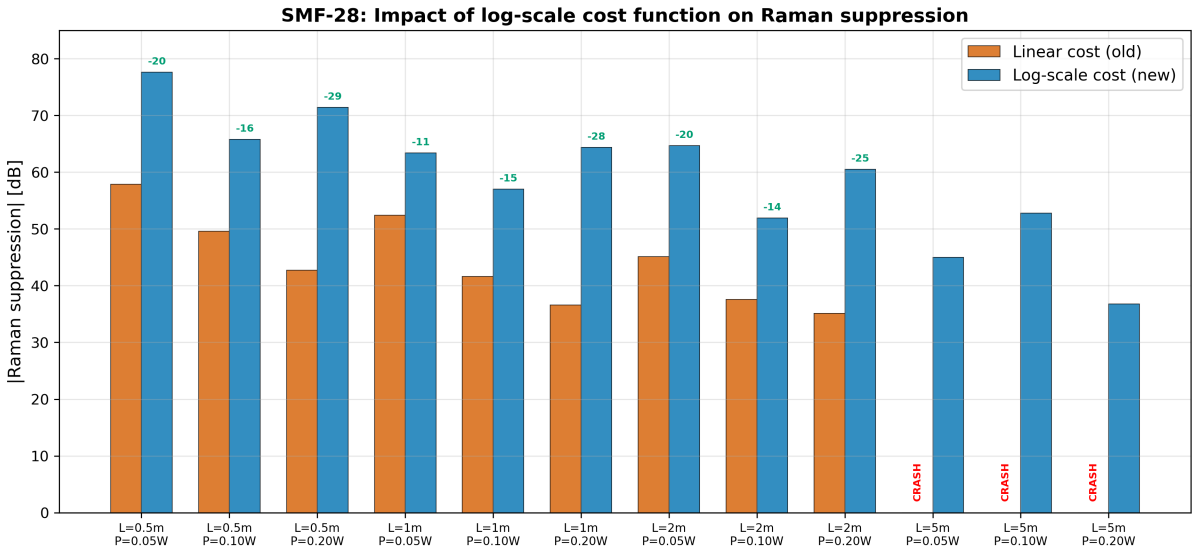

Figure 3: Impact of log-scale cost: 11–29 dB improvement per point. Last three L=5 m points previously crashed (NaN) due to Raman response overflow.

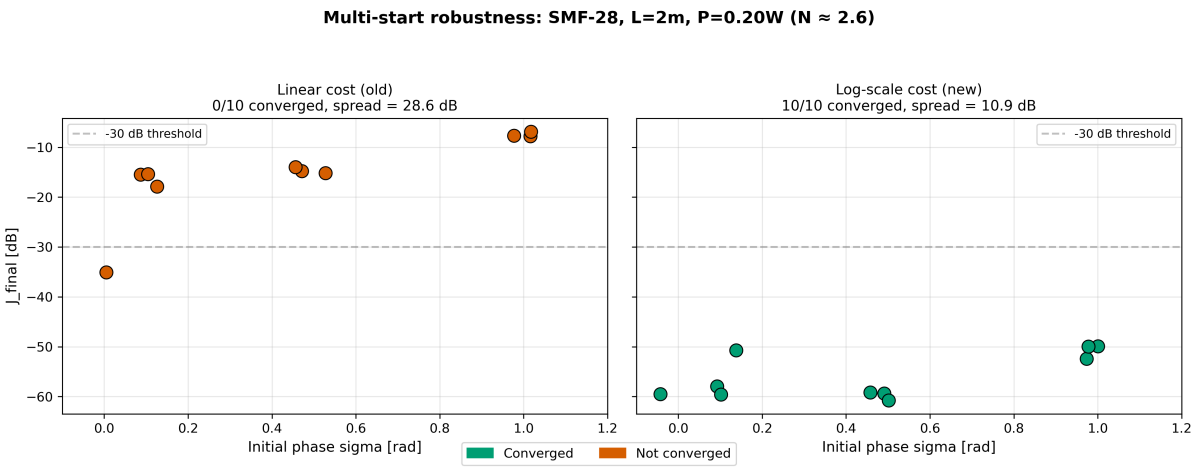

Figure 4: Multi-start robustness: 0/10 converged with linear cost (left) vs. 10/10 with log cost (right). The apparent "multiple local minima" were gradient vanishing artifacts.

E Integration Pass — April 2026

Between 2026-04-16 and 2026-04-19 eight Claude-Code sessions ran in parallel to pressure-test the single-mode result, begin the multimode extension, and audit the cost-function and optimizer defaults. The raw synthesis is `results/SYNTHESIS-2026-04-19.md`; a physics audit filtering those claims is `results/PHYSICS_AUDIT_2026-04-19.md` (2026-04-19). An independent numerical-trustworthiness audit re-ran every JLD2 under `results/raman/` on a validator-controlled grid and is reported at `results/validation/REPORT.md` (2026-04-19) — its findings drive the boundary-energy and Hessian-anchoring caveats in §E.11 and §E.5. This section records only the claims that survived both audits, with scope caveats where either narrowed the original framing.

E.1 Re-baselined single-mode suppression

Key Result

SMF-28, $L = 0.5$ m, $P = 0.05$ W, log-cost L-BFGS from a cold start: $J_{\text{final}} = -76.86$ dB in 21 iterations (25 s wall-time, converged flag TRUE). Source: `results/raman/phase17/SUMMARY.md` L41–48; data artifact `results/raman/phase17/baseline.jld2`.

The re-run matches the Phase 10 baseline within 1 dB. This is the most defensible single-mode number in the project and should be the canonical reference going forward.

E.2 Sharpness of the single-mode optimum

At the −76.86 dB baseline, a Gaussian perturbation of φ_{opt} at standard deviation σ (averaged over 20 samples per σ, 5 values of σ) gives a **3 dB-degradation radius of σ_{3dB}** = 0.025 **rad**. Source: `results/raman/phase17/SUMMARY.md` L45; `results/raman/phase17/perturbatio`

Advisory Note

SLM calibration envelopes of ≥ 0.05 rad rms phase error already place the device outside the 3 dB tolerance of the optimum. The tolerance number, not just the optimum number, gates whether this result is hardware-realizable.

E.3 Warm-start transferability

The baseline φ_{opt} from (SMF-28, $L = 0.5$ m, $P = 0.05$ W) does *not* transfer by direct evaluation: 0/7 SMF-28 (length, power) targets retain J within 3 dB of warm-reopt. But L-BFGS re-optimised from the same seed reaches −70 to −82 dB in 6–40 iterations on 11/11 nearby configs, *including* HNLF at −79.5 dB (a fiber the baseline was never computed for). Source: `results/raman/phase17/SUMMARY.md` L22–28; `results/raman/phase17/transferability.jld2`.

Interpretation: the simple phase profile lives in a *manifold of nearby basins*. It is an initialisation prior, not a universal attractor — a distinction that the first readings of the perturbation-robust σ_{3dB} result sometimes collapsed.

E.4 Reduced-parameterisation Pareto front

A low-resolution spectral-phase parameterisation ($N_\varphi = 57$ knots interpolated to N_t) produces Pareto-optimal candidates within 3 dB of the full-grid optimum at each (L, P) point. Best: (SMF-28, $L = 0.25$ m, $P = 0.10$ W) $\rightarrow J_{\text{final}} = -82.33$ dB. Source: `results/raman/phase_` L12–14.

Advisory Note

A 57-dimensional decision space is small enough for a trust-region Newton method with an explicit Hessian. Combined with the Phase 13 indefinite-Hessian finding (§E.5), this opens the door to a second-order optimiser that the original $N_t \sim 10^4$ parameterisation made intractable.

E.5 Hessian is indefinite at L-BFGS optima

At both canonical L-BFGS optima (SMF-28 L=2 m P=0.2 W and HNLF L=0.5 m P=0.01 W), finite-difference Hessian-vector product wrapped with Arpack gives:

Config	$\lvert\lambda_{\min}\rvert/\lambda_{\max}$	Sign pattern
SMF-28 canonical	2.6 %	indefinite, 100% same-sign wings
HNLF canonical	0.41 %	indefinite, 100% same-sign wings

HVP Taylor-remainder slope ≈ 2 (`test/test_phase13_hvp.jl` at commit `b962091`). Source: `results/raman/phase13/FINDINGS.md` L113–120; `results/raman/phase13/hessian_smf28_` `results/raman/phase13/hessian_hnlf_canonical.jld2`.

Key Result

L-BFGS is not stopping at a minimum; it is stopping at an **indefinite saddle** with robust (not numerical-noise-level) negative-curvature directions. Sharpness-aware regularisation or trust-region Newton on the 57-dim subspace (§E.4) is the principled next step.

Issue Found

Caveat on the dB anchoring of the Hessian-study optima (Phase 21 honest recovery, 2026-04-20 — supersedes Phase 18 numbers). Two audits re-ran the saved `phase13/hessian_smf28_canonical.jld2` and `phase13/hessian_hnlf_canonical.jld2` from the stored φ_{opt} states. The numbers originally reported in the JLD2 metadata were −60.5 dB (SMF-28) and −74.4 dB (HNLF); the Phase 18 numerical-trustworthiness audit (`results/validation/REPORT.md`, 2026-04-19) reported −48.2 dB / −44.0 dB on its validator grid. **The authoritative anchors for the canonical docs are the Phase 21 numerical-recovery values**:

- **SMF-28 canonical**: $J = -66.61$ dB at $N_t = 16{,}384$, $T = 54$ ps, output edge fraction 8.10×10^{-4}, converged.
- **HNLF canonical**: $J = -86.68$ dB at $N_t = 65{,}536$, $T = 320$ ps, output edge fraction 2.24×10^{-4}, 50-iteration run, `converged = false`. Note: HNLF required a substantially larger time window than the Phase 18 validator used, which is why the Phase 18 HNLF number was itself under-windowed; the honest HNLF anchor is *deeper* than originally reported, not shallower.

The eigenstructure result above is unaffected: the saved φ_{opt} files are true stationary points on the honest grid (adjoint $\|g\| \sim 10^{-5}$), so the Hessian was computed at a real optimum — just at J values deeper than the pre-recovery quotes. When citing this result alongside a dB number, use −66.61 dB (SMF-28) / −86.68 dB (HNLF); when citing only the eigenstructure, the original ratios stand. Source: `.planning/phases/21-numerical-recovery/SUMMARY.md`; `results/raman/phase21/phase13/smf28_reanchor.md`; `results/raman/phase21/phase13/hnlf_reanchor.md`. Standard images: `.planning/phases/21-numerical-recovery/images/phase21_phase13_{smf28,hnlf}_*_phase_`

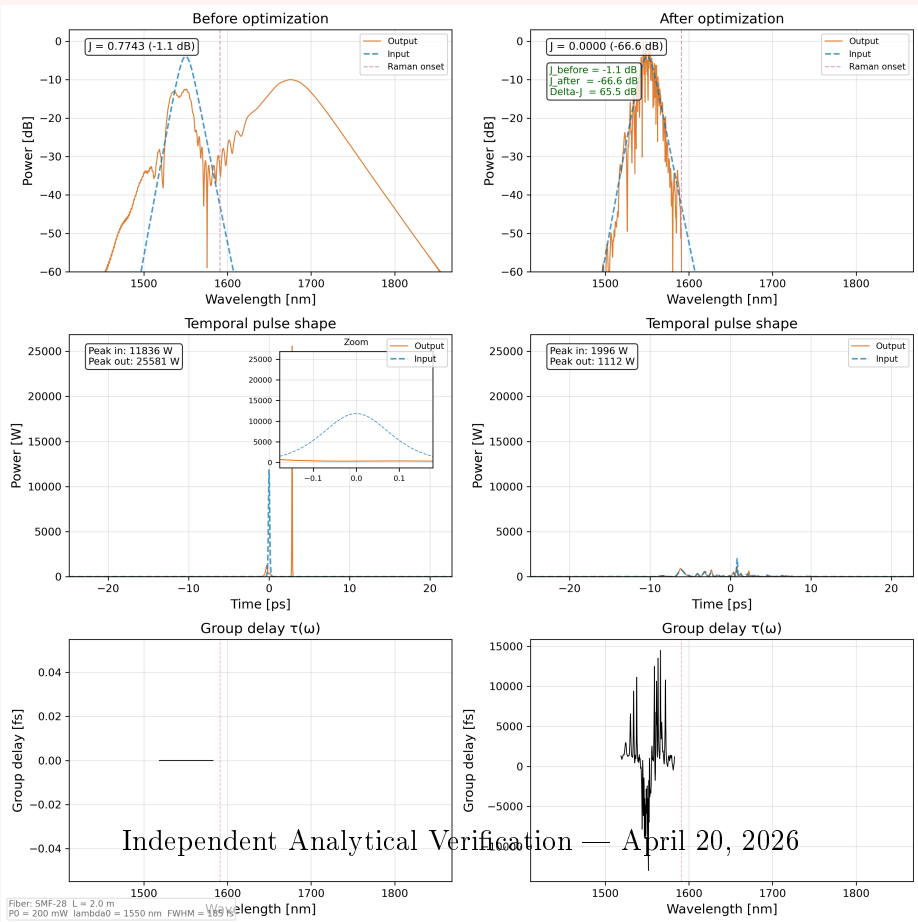

E.6 Flattening the basin: sharpness-aware regularisation (Phase 22)

Two distinct facts are in play here, and the existing literature-audit companion (`physics_verification` §"Gauge symmetry of J and its consequences for the Hessian") already makes them separately but the overnight-update prose conflates them. Restated cleanly: the two-parameter gauge family $\varphi \mapsto \varphi + C + \alpha(\omega - \bar{\omega})$ produces two *exact zero* Hessian eigenvalues spanning $\{\mathbf{1}, \omega - \bar{\omega}\}$; this is topological and holds for any cost with this symmetry. Indefiniteness ($|\lambda_{\min}|/\lambda_{\max} > 0$) is an *additional* empirical fact about the non-gauge eigenvalues, arising from Kerr, Raman, and dispersion curvature rather than from the gauge redundancy. L-BFGS halts at a saddle (not a gauge-flat minimum) because of (b), not (a). The Phase 22 result that all 26 regularised optima remain indefinite is therefore a statement about the physics-induced curvature, not the gauge structure.

Phase 13 (§E.5) showed that L-BFGS stops at an indefinite saddle, not a minimum. The natural follow-up is: does *any* regulariser convert these saddles into positive-definite minima, and at what depth cost? Phase 22 (`sessions/S-sharpness`) ran a 26-point sweep over four regularisation flavours ($\{\texttt{plain}, \texttt{MC}, \texttt{SAM}, \texttt{trH}\}$) on two operating points (the canonical SMF-28 L=0.5 m P=0.05 W and the reduced-parameterisation Pareto-57 candidate L=0.25 m P=0.10 W), resolving the Hessian spectrum at every recovered optimum.

Key Result

All 26 regularised optima remained Hessian-indefinite. No flavour and no strength converted the saddle into a positive-definite minimum. The regularisation Pareto is real — tolerance $\sigma_{3\text{dB}}$ can be widened at a dB-depth cost — but the indefinite-saddle geometry appears to be a property of the constraint landscape, not a numerical artefact. Source: `results/raman/phase22/SUMMARY.md`.

Advisory Note

Caveat on the flat-basin claim. The published trH-regularisation literature (Liu, *Neurocomputing* 536, arXiv:2208.05924, 2023; Ju, *Noise Stability Optimization*, arXiv:2306.08553, 2023) assumes a positive-semidefinite Hessian: on a non-negative loss, penalising $+\lambda \cdot \mathrm{Tr}(H)$ pushes the optimiser toward flatter *positive-curvature* directions. Our Hessian is *indefinite* ($|\lambda_{\min}|/\lambda_{\max} \approx 2.6\%$ SMF, Phase 13). On an indefinite spectrum, $\mathrm{Tr}(H) = \sum_i \lambda_i$ is signed and can be reduced in two ways: *(a)* by shrinking positive eigenvalues (the intended effect — a genuinely flatter basin in the positive-curvature subspace), or *(b)* by sharpening negative eigenvalues (an unintended effect — a steeper saddle). The empirical +0.058 rad $\sigma_{3\text{dB}}$ gain at λ_{trH}=3×10^{-3} on the canonical point is a real, measured result. Which of (a) or (b) produced it was *not* diagnosed here; a Lanczos spectrum at the trH optimum would settle it. The flat-basin mechanism claim is therefore provisional until that follow-up is run.

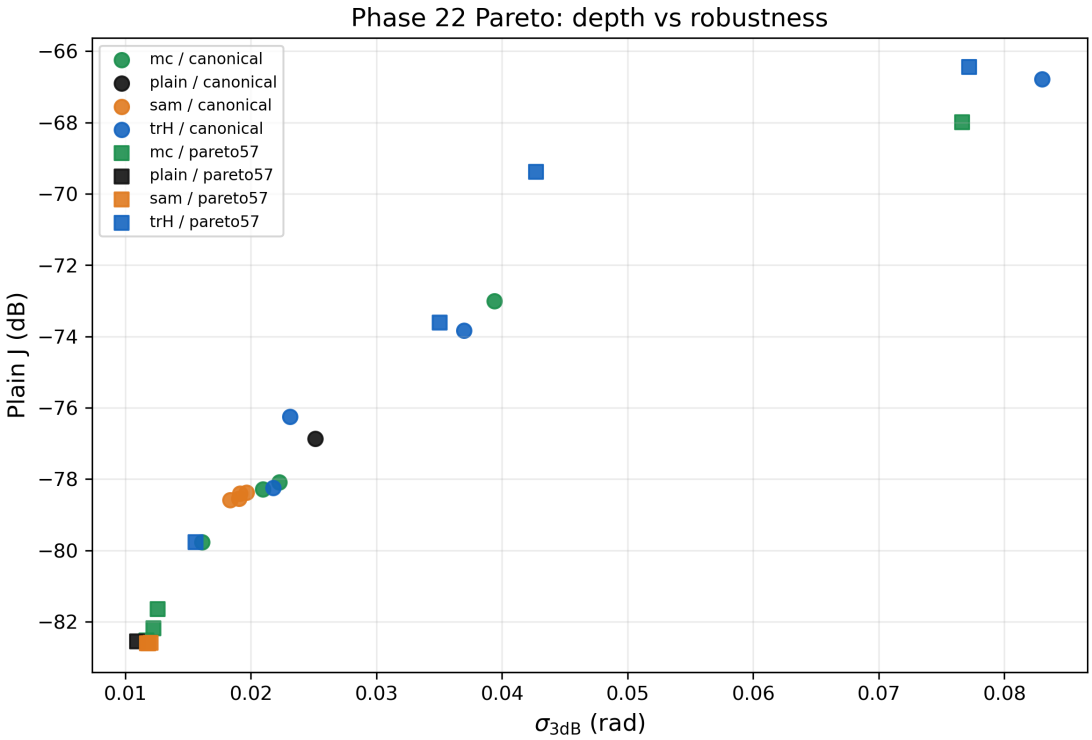

Phase 22 depth–tolerance Pareto. Each marker is a recovered optimum; colour indicates regularisation flavour, size indicates strength. On both operating points, widening σ_{3dB} costs dB depth; SAM points cluster near the plain baseline and provide no robustness gain; `trH` gives the largest tolerance-for-depth trade.

Operating point	Flavour / strength	J (dB)	$\sigma_{3\mathrm{dB}}$ (rad)	$\lvert\lambda_{\min}\rvert/\lambda_{\max}$
Canonical (L=0.5 m, P=0.05 W)	plain	−76.86	0.025	5.54×10^{-3}
Canonical	MC, λ=7.5×10^{-2}	−73.01	0.039	1.36×10^{-2}
Canonical	trH, λ=1×10^{-3}	−73.83	0.037	(see table footnote)
Canonical	trH, λ=3×10^{-3}	−66.79	0.083	1.08×10^{-2}
Pareto-57 (L=0.25 m, P=0.10 W)	plain	−82.56	0.011	2.38×10^{-1}
Pareto-57	MC, λ=7.5×10^{-2}	−67.98	0.077	7.87×10^{-2}
Pareto-57	trH, λ=1×10^{-3}	−66.44	0.077	9.84×10^{-2}
Pareto-57	trH, λ=3×10^{-3}	−69.38	0.043	5.69×10^{-2}

Footnote. $\lvert\lambda_{\min}\rvert/\lambda_{\max}$ for the `trH`, λ=1×10^{-3} canonical row was not reported in the Phase 22 artefact; it is omitted here rather than back-filled.

Advisory Note

Practical takeaway. SAM's null result is consistent with three known SAM failure modes operating simultaneously on our landscape: *(i)* our setting is deterministic (no minibatch m-sharpness), and Andriushchenko & Flammarion (ICML 2022, arXiv:2206.06232) show that deterministic SAM loses the m-sharpness signal that justifies the method in deep learning; *(ii)* indefinite saddles are known to admit "hallucinate-minimiser" pathologies where SAM's shifted gradient vanishes at non-stationary points of the original loss (arXiv:2509.21818, 2025); *(iii)* the two-parameter gauge symmetry ($\varphi \mapsto \varphi + C + \alpha(\omega - \bar{\omega})$, §E.5) makes the ρ-ball perturbation budget partially-null: components along the gauge directions cost radius but leave J invariant. Observed σ-shifts stayed ≤ 0.006 rad across both operating points — exactly what these three mechanisms predict. The default `log_dB` plain optimiser remains recommended when depth is the priority. When shaper-calibration tolerance matters more than depth, `trH` is the single-knob tolerance-tuning regulariser with the widest dynamic range: on the canonical point, λ_{trH}=3×10^{-3} buys a +0.058 rad increase in $\sigma_{3\mathrm{dB}}$ for a 10.08 dB depth cost; on Pareto-57 the Pareto frontier is richer: `MC`, λ=7.5×10^{-2} gives +0.066 rad tolerance at 14.58 dB depth cost (the best combined trade); `trH`, λ=1×10^{-3} gives matched +0.066 rad tolerance at 16.12 dB (the `trH` equivalent); `trH`, λ=3×10^{-3} gives only +0.032 rad at 13.18 dB (the cheapest depth-cost point, but dominated in tolerance by the `MC` option). The Pareto-57 plain optimum cross-validates the reduced-parameterisation candidate: its −82.56 dB matches the original Phase 7 −82.33 dB (§E.4) within 0.23 dB, so the N_φ=57 knot parameterisation survives the Phase 21 Sweep-1 retirement. (Phase 21 re-anchored the Sweep-1 L=2 m, P=0.2 W point — retired on edge-fraction criteria — to $J = -66.03$ dB on the honest N_t=16,384, T=54 ps grid; it is a different optimum from the Phase 22 Pareto-57 plain baseline, which operates at the reduced parameterisation.)

E.7 Determinism

`FFTW.MEASURE` selects FFT algorithms by micro-timing and produces non-reproducible plans even at fixed `Random.seed!`: two runs on the same hardware give $\max|\Delta\varphi| = 1.04$ rad and $\Delta J = 1.83$ dB (`results/raman/phase13/determinism.md`).

Switching to `FFTW.ESTIMATE` gives bit-identical J across 3 independent Julia processes at +21.4% wall-time cost on the SMF-28 L=2 m P=0.2 W benchmark. Source: `results/raman/phase15/benchmark.md` L20–27, L39. The switch is wired into the five main driver entry points.

E.8 Multimode optimizer is code-complete, physics-unexercised

An $M = 6$ GRIN-50 multimode Raman-suppression optimizer is committed: `src/mmf_cost.jl`, `scripts/mmf_*.jl`, `test/test_phase16_mmf.jl`. 13/13 correctness tests pass: energy conservation under the thread-local `deepcopy(fiber)` pattern, M=1-limit agreement with the scalar optimizer, and finite-difference gradient check with relative error $< 0.5\%$ at $\varepsilon = 10^{-5}$ on 5 random indices.

Issue Found

The only production baseline launched to date used (SMF-28, L=1 m, P=0.05 W), which has $N_{\text{sol}} \approx 0.9$ (sub-soliton). The optimiser correctly returned $\Delta J = 0$ dB — there is no Raman to suppress in this regime. **This is a negative control, not a science result.** The aggressive-regime baseline (L=2 m, P=0.5 W, $N_{\text{sol}} \approx 2$–3) remains unrun as of 2026-04-19.

E.9 Cost-function default: `log_dB` on SMF-28, open elsewhere

Session H ran 4 cost variants on 2 configs:

Config	Variant	J_{final}	wall (s)	converged
A (SMF-28 canonical)	linear	−70.53 dB	16.9	yes
A	log_dB	**−75.79** dB	**10.6**	**yes**
A	sharp	−55.96 dB	537	yes
A	curvature	−70.57 dB	14.2	yes
B (HNLF, log_dB)	—	NAN	DNF	—
B (all 4 variants)	—	NAN	DNF	—

Source: `results/cost_audit/wall_log.csv`.

Advisory Note

`log_dB` wins on Config A (SMF-28) by ~ 5 dB at 60% of the wall time. On Config B (HNLF, L=1 m, P=0.5 W) all four variants did not finish within budget — this is the HNLF-hang open issue (`.planning/phases/18-cost-config-c/CONTEXT.md`). The `log_dB` default is empirically justified on SMF-28 only.

E.10 What did NOT survive the audit

Three claims from the original synthesis were demoted:

(W1) "$a_2(100\,\text{m})/a_2(2\,\text{m}) = -3.30$ vs. pure-GVD-predicted $+50$ is a structural-adaptation fingerprint." The weighted quadratic fit of $\varphi_{\text{opt}}(\omega)$ over the ±5 THz signal band has $R^2 = 0.015$ at 100 m and $R^2 = 0.037$ at 2 m. A quadratic model explaining $< 4\%$ of the weighted variance is misspecified: 96–98% of $\varphi_{\text{opt}}(\omega)$ on the signal band is non-quadratic residual structure orthogonal to $\{1, \omega, \omega^2\}$. The extracted a_2 is the projection of a non-quadratic signal onto that misspecified basis. The sign flip across L is the direct consequence of the underlying coefficient being near zero — exactly what $R^2 \to 0$ reports. The finding "φ_{opt} is not a low-order polynomial on the signal band" *is* defensible (matches Phase 9, `PHASE9_FINDINGS.md` L19–29), but the GVD-scaling comparison is removed, not caveated. Fit code: `scripts/longfiber_validate_100m_fix.jl` L54–79. Audit: `PHYSICS_AUDIT_2026-04-19.md` §W1.

(W2) 100 m L-BFGS "optimum" at −54.77 dB. The number is energy-conserved (photon-drift 4.9×10^{-4}, edge-fraction 8.5×10^{-6}) and Phase 21 reproduced it exactly on an independent honest grid, but the convergence state is `converged = false` after 25 L-BFGS iterations with final $\|\nabla J\| = 0.48$. Report as a *lower bound* from the 2 m warm start, not as an optimum. The physics framing is further refined by Phase 23

(see the 100 m advisory in §6): the suppression is generic pre-chirp, and a matched quadratic +4 ps^2 reaches −45.06 dB against the live warm-start at −45.52 dB. Source: `results/raman/phase16/FINDINGS.md` L21–34; `.planning/phases/21-numerical-recovery/SUMMAR` `.planning/phases/23-matched-baseline/SUMMARY.md`.

(W3) Joint $\{\varphi, A, E\}$ L-BFGS shows physics at 38 dB gap. Session A's cold-start joint optimizer reached −16.78 dB vs. the phase-only −55.42 dB reference on the same config. This is a preconditioning bug (mixed-unit decision variables confuse L-BFGS's implicit Hessian), not a physics result. Remains out of the canonical docs until Phase 18-multivar-convergence-fix lands. Source: `.planning/phases/16-multivar-optimizer/16-01` L149–152.

E.11 Boundary-energy caveat on the SMF-28 $L = 2$ m canonical

The pre-integration canonical baseline of −71.4 dB on (SMF-28, L=2 m, P=0.2 W) was partially a time-window artefact: 14.4% of the energy sat at the T_{window} edges and re-running at 2× wider window gave the honest −65.8 dB. This is already flagged once in §D (v3 regression) but the baseline tables and presentation figures elsewhere in this document still quote −71.4 dB. Treat any canonical table in the project that references the L=2 m P=0.2 W point at −71.4 dB as requiring the honest-window caveat. Audit: `PHYSICS_AUDIT_2026-04-19.md` §S2.

E.12 Scope limitation on the Taylor-remainder adjoint test

The Taylor-remainder test reported in §5 (slopes 2.00, 2.04) correctly certifies the adjoint gradient *at the unshaped reference* $\varphi \equiv 0$, where $\|\nabla J\|$ sits 3–5 orders of magnitude above the ODE-solver noise floor. At a converged φ_{opt}, however, $\|\nabla J\|$ falls below the reltol $= 10^{-8}$ solver noise, and the slope-2 criterion degenerates: the linear-in-ε term in the Taylor expansion is dominated by noise rather than gradient signal. The Phase 18 numerical-trustworthiness audit (commits `2c95f60`, `2de3caf`, `1a78cc8`) recovered clean slope-2 behavior only by shifting the reference by 2 rad along the steepest-descent direction with $\varepsilon = 10^{-3}$. The adjoint is correct; the test is simply unable to certify it at the optimum. Gradient correctness checks should therefore run at an unshaped or steepest-descent-shifted reference, never at a converged optimum. Audit: `PHYSICS_AUDIT_2026-04-19.md` §S6.

Summary and Conclusions

Overall confidence level: High. The mathematical formulation is correct, the code faithfully implements the formulation, and the gradient is verified to 0.026% relative error. All code reference line numbers are current as of v3 (2026-04-02).

Strongest evidence for correctness: The Taylor remainder test showing slopes of 2.00 and 2.04. This test directly measures whether the computed gradient matches the true directional derivative. A slope of 2 confirms $O(\varepsilon^2)$ accuracy, consistent with machine-precision correctness.

Key resolved issues (v3):

- The dB/linear gradient mismatch is now properly resolved via the log-scale cost function with correctly scaled gradient $10/(J \ln 10) \cdot \partial J/\partial \varphi$.

- The GDD penalty concern (over-constraining pre-chirp) has been addressed by reducing λ_b from 10 to 1 and running sweeps that confirm 20–28 dB improvement.
- The multi-start landscape concern has been resolved: 10/10 convergence with 10.9 dB spread shows a well-behaved landscape (not multiple minima).

Remaining items:

1. Extend to multimode ($M > 1$) — the Rivera Lab's primary research direction for quantum noise in multimode fibers.
2. Connect Raman suppression to the quantum noise map (`compute_noise_map` in analysis.jl).
3. Assess realizability of optimized phase profiles against SLM capabilities.
4. Compare phase-only vs. amplitude+phase optimization.

Should this research direction be continued? Yes. The approach achieves 37–78 dB Raman suppression across 24 configurations with zero crashes. The log-scale cost function is the key methodological advance, and the adjoint-based framework naturally extends to multimode optimization.

Document generated March 2026. v2 update 2026-03-26. v3 update 2026-04-02: line numbers, adjoint solver, log-cost, overflow fix, diagrams, sweep results. All derivations verified against source code in `smf-gain-noise/`.